

NNDL_Project_Notebook_Live_Submit

December 16, 2025

1 SECTION 1: Setup & Imports

```
[ ]: from google.colab import drive
drive.mount('/content/drive')

!nvidia-smi # just to sanity-check the GPU

!pip install wandb -q
import wandb

USE_WANDB = True
WANDB_ENTITY = "nndl-project-F25"
WANDB_PROJECT = "Multihead-Classification-Competition"

if USE_WANDB:
    wandb.login()

from datetime import datetime
import os, zipfile, random

import numpy as np
import pandas as pd
from PIL import Image

import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim

from torch.utils.data import Dataset, DataLoader, random_split
from torchvision import transforms, models

import matplotlib.pyplot as plt

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print("Using device:", device)
```

```

# wandb run-naming schema
def make_run_name(base: str) -> str:
    """Create a unique run name with timestamp."""
    return f'{base}_{datetime.now().strftime('%Y%m%d_%H%M%S')}'"

# optional: for approximate reproducibility
def set_seed(seed=42):
    random.seed(seed)
    np.random.seed(seed)
    torch.manual_seed(seed)
    torch.cuda.manual_seed_all(seed)

set_seed(42)

```

Drive already mounted at /content/drive; to attempt to forcibly remount, call
drive.mount("/content/drive", force_remount=True).

Sun Dec 14 18:21:33 2025

```

+-----+
-----+
| NVIDIA-SMI 550.54.15           Driver Version: 550.54.15      CUDA Version:
12.4      |
|-----+-----+-----+
-----+
| GPU  Name                  Persistence-M | Bus-Id          Disp.A | Volatile
Uncorr. ECC |
| Fan  Temp     Perf            Pwr:Usage/Cap |          Memory-Usage | GPU-Util
Compute M. |
|          |                               |          |
MIG M.  |
|=====+=====+=====+=====+=====+=====+
=====|
|   0  NVIDIA A100-SXM4-80GB        Off | 00000000:00:05.0 Off |
0 |
| N/A  32C     P0                52W / 400W | 0MiB / 81920MiB | 0%
Default |
|          |                               |          |
Disabled |
+-----+-----+-----+
-----+
-----+
| Processes:
|
| GPU  GI  CI          PID  Type  Process name
GPU Memory |

```

```

|      ID    ID
Usage   |
=====
=====
| No running processes found
|
+-----+
-----+
wandb: Currently logged in as: akseldkw
(akseldkw07) to https://api.wandb.ai. Use `wandb login
--relogin` to force relogin
Using device: cuda

```

2 SECTION 2: Config & Paths

```
[ ]: DATA_ROOT = "/content/drive/MyDrive/NNDL-Project/Project Data"

# Local scratch space on the VM / Colab
LOCAL_DATA_ROOT = "/content/local_data"
os.makedirs(LOCAL_DATA_ROOT, exist_ok=True)

train_zip_path = os.path.join(DATA_ROOT, "train_images.zip")
test_zip_path = os.path.join(DATA_ROOT, "test_images.zip")

# Unzip to LOCAL_DATA_ROOT instead of Drive
train_out_dir = os.path.join(LOCAL_DATA_ROOT, "train_images")
test_out_dir = os.path.join(LOCAL_DATA_ROOT, "test_images")

if not os.path.exists(train_out_dir):
    os.makedirs(train_out_dir, exist_ok=True)
    with zipfile.ZipFile(train_zip_path, 'r') as z:
        z.extractall(train_out_dir)

if not os.path.exists(test_out_dir):
    os.makedirs(test_out_dir, exist_ok=True)
    with zipfile.ZipFile(test_zip_path, 'r') as z:
        z.extractall(test_out_dir)

TRAIN_IMG_DIR = os.path.join(train_out_dir, "train_images")
TEST_IMG_DIR = os.path.join(test_out_dir, "test_images")

TRAIN_CSV = os.path.join(DATA_ROOT, "train_data.csv")
SUPER_CSV = os.path.join(DATA_ROOT, "superclass_mapping.csv")
SUB_CSV = os.path.join(DATA_ROOT, "subclass_mapping.csv")
```

```

BATCH_SIZE = 64
NUM_WORKERS = 2 # can set to 0 if we hit dataloader issues

VAL_SPLIT = 0.1 # 10% validation
IMG_SIZE = 64 # our image dimensions

PROJECT_NAME = "coms4776-transfer-learning" # TBD update
APPROACH = "two_models" # "two_heads" or "two_models"
DATA_AUGMENT = True

# Indices for "novel" classes (per provided data)
NOVEL_SUPER_IDX = 3 # superclass index for novel
NOVEL_SUB_IDX = 87 # subclass index for novel

# Number of times run full-batch
EPOCHS = 15

# Learning rates
LR = 1e-4 # overall learning rate
LR_HEAD = 1e-2 # head learning rate, used when freezing backbone
WEIGHT_DECAY = 1e-4 # seems standard
BACKBONE = "resnet50" # "resnet18" or "resnet50"

# Novel-super CIFAR integration (more images)
# Options: "none", "small" (~1000 samples), "large" (~5000 samples)
CIFAR_NOVEL_MODE = "large" # "large" or "small" or "none"

# Path to store metadata about CIFAR novel images
CIFAR_NOVEL_CSV_PATH = os.path.join(LOCAL_DATA_ROOT, "cifar_novel_data.csv")

# Fine-tuning mode for ResNet backbone
# "full" = train all layers (what you're currently doing)
# "frozen" = freeze backbone, train only the heads on top
FINE_TUNE_MODE = "full" # or "frozen"

# Initial novelty thresholds (starting points, will tune further)
TAU_SUPER = 0.99 # NOTE: per calibration with validation data. if max_ ↵superclass prob < TAU_SUPER -> predict novel superclass
TAU_SUB = 0.85 # NOTE: per calibration with validation data. if max_ ↵subclass prob < TAU_SUB -> predict novel subclass

##### MAKE SURE USE_PSEUDO_NOVEL IS FALSE BEFORE LEADERBOARD SUBMISSION #####
# Used to fine-tune TAU_SUB
USE_PSEUDO_NOVEL = True # to validate on held-out subclasses from training. ↵Used to fine-tune TAU_SUB
PSEUDO_NOVEL_FRACTION = 0.15
PSEUDO_NOVEL_SEED = 123

```

```
[ ]: # Build CIFAR-100 novel-super dataset (excluding reptiles)
from torchvision.datasets import CIFAR100

# Download CIFAR100 once (raw PIL images)
CIFAR_ROOT = os.path.join(LOCAL_DATA_ROOT, "cifar100_raw")
os.makedirs(CIFAR_ROOT, exist_ok=True)

# Only do the heavy image-copying if CSV doesn't exist
if CIFAR_NOVEL_MODE != "none" and not os.path.exists(CIFAR_NOVEL_CSV_PATH):
    print("Building CIFAR novel-super dataset (this happens once)...")


    cifar_train = CIFAR100(root=CIFAR_ROOT, train=True, download=True, ↴
    transform=None)

    # CIFAR-100 fine label names (with underscores)
    cifar_fine_names = cifar_train.classes # e.g. "apple", "beaver", ↴
    "aquarium_fish", ...

    # Fine classes we want, based on your list, excluding reptiles
    # (these names match CIFAR-100 fine label names)
    allowed_fine_names = set([
        # aquatic mammals
        "beaver", "dolphin", "otter", "seal", "whale",
        # fish
        "aquarium_fish", "flatfish", "ray", "shark", "trout",
        # flowers
        "orchid", "poppy", "rose", "sunflower", "tulip",
        # food containers
        "bottle", "bowl", "can", "cup", "plate",
        # fruit and vegetables
        "apple", "mushroom", "orange", "pear", "sweet_pepper",
        # household electrical devices
        "clock", "keyboard", "lamp", "telephone", "television",
        # household furniture
        "bed", "chair", "couch", "table", "wardrobe",
        # insects
        "bee", "beetle", "butterfly", "caterpillar", "cockroach",
        # large carnivores
        "bear", "leopard", "lion", "tiger", "wolf",
        # large man-made outdoor things
        "bridge", "castle", "house", "road", "skyscraper",
        # large natural outdoor scenes
        "cloud", "forest", "mountain", "plain", "sea",
        # large omnivores and herbivores
        "camel", "cattle", "chimpanzee", "elephant", "kangaroo",
        # medium-sized mammals
        "fox", "porcupine", "possum", "raccoon", "skunk",
    ])
```

```

# non-insect invertebrates
"crab", "lobster", "snail", "spider", "worm",
# people
"baby", "boy", "girl", "man", "woman",
# small mammals
"hamster", "mouse", "rabbit", "shrew", "squirrel",
# trees
"maple", "oak", "palm", "pine", "willow",
# vehicles 1
"bicycle", "bus", "motorcycle", "pickup_truck", "train",
# vehicles 2
"lawn_mower", "rocket", "streetcar", "tank", "tractor",
# NOTE: reptiles group ("crocodile", "dinosaur", "lizard", "snake", "turtle") is *excluded* on purpose
])

# Map class_name -> list of indices for that fine class
name_to_indices = {name: [] for name in allowed_fine_names}
for idx in range(len(cifar_train)):
    _, fine_label = cifar_train[idx] # fine_label is int index into cifar_fine_names
    fine_name = cifar_fine_names[fine_label]
    if fine_name in allowed_fine_names:
        name_to_indices[fine_name].append(idx)

# Flatten candidate indices across all allowed classes
candidate_indices = []
for name, idx_list in name_to_indices.items():
    candidate_indices.extend(idx_list)

print("Total candidate CIFAR images (allowed classes, excl. reptiles):", len(candidate_indices))

# Target total novel-super samples (max 5000, as you requested)
TARGET_TOTAL = 5000
random.seed(42)
random.shuffle(candidate_indices)
selected_indices = candidate_indices[:TARGET_TOTAL]

print("Selected indices:", len(selected_indices))

# Copy images into TRAIN_IMG_DIR and build CSV rows
cifar_aug_records = []

for idx in selected_indices:
    img, fine_label = cifar_train[idx]
    fine_name = cifar_fine_names[fine_label]

```

```

# Unique filename to avoid collisions
fname = f"cifar_novel_{idx}_{fine_name}.png"
dst_path = os.path.join(TRAIN_IMG_DIR, fname)

# img is PIL.Image when transform=None
img.save(dst_path)

record = {
    "image": fname,
    "superclass_index": NOVEL_SUPER_IDX, # novel superclass
    "subclass_index": NOVEL_SUB_IDX, # novel subclass
    "description": f"CIFAR100:{fine_name} (novel superclass)",
}
cifar_aug_records.append(record)

cifar_aug_df = pd.DataFrame(cifar_aug_records)
cifar_aug_df.to_csv(CIFAR_NOVEL_CSV_PATH, index=False)
print("Saved CIFAR novel-super metadata to:", CIFAR_NOVEL_CSV_PATH)
print("Images copied into TRAIN_IMG_DIR:", TRAIN_IMG_DIR)

elif CIFAR_NOVEL_MODE != "none":
    print("CIFAR novel-super CSV already exists at:", CIFAR_NOVEL_CSV_PATH)
else:
    print("CIFAR_NOVEL_MODE='none'; skipping CIFAR novel-super creation.")

```

CIFAR novel-super CSV already exists at:
 /content/local_data/cifar_novel_data.csv

3 SECTION 3: Data Loading & DataLoaders

```
[ ]: # SECTION 3: Data Loading & Dataloaders

# Base training data from class
base_train_df = pd.read_csv(TRAIN_CSV)

super_map_df = pd.read_csv(SUPER_CSV) # columns: index, class
sub_map_df = pd.read_csv(SUB_CSV) # columns: index, class

num_super = len(super_map_df)
num_sub = len(sub_map_df)

print("Num superclasses:", num_super)
print("Num subclasses:", num_sub)

# --- Integrate CIFAR novel-super examples, if enabled ---
```

```

if CIFAR_NOVEL_MODE == "none":
    train_df = base_train_df.copy()
    print("CIFAR_NOVEL_MODE='none' → using only original training data.")
else:
    if not os.path.exists(CIFAR_NOVEL_CSV_PATH):
        raise FileNotFoundError(
            f"CIFAR_NOVEL_MODE={CIFAR_NOVEL_MODE} but {CIFAR_NOVEL_CSV_PATH} was not found. "
            "Run the CIFAR novel-super build section first."
        )

cifar_aug_df = pd.read_csv(CIFAR_NOVEL_CSV_PATH)

if CIFAR_NOVEL_MODE == "small":
    # Use ~1000 CIFAR novel-super images
    n_small = min(1000, len(cifar_aug_df))
    cifar_aug_df = cifar_aug_df.sample(n=n_small, random_state=42).
    reset_index(drop=True)
    print(f"CIFAR_NOVEL_MODE='small' → using {len(cifar_aug_df)} CIFAR novel-super samples.")

elif CIFAR_NOVEL_MODE == "large":
    # Use all available (up to 5000 created earlier)
    print(f"CIFAR_NOVEL_MODE='large' → using {len(cifar_aug_df)} CIFAR novel-super samples.")

else:
    raise ValueError(f"Unknown CIFAR_NOVEL_MODE: {CIFAR_NOVEL_MODE}")

# Combine original training data with CIFAR novel-super rows
train_df = pd.concat([base_train_df, cifar_aug_df], ignore_index=True)
print("Combined train_df size (original + CIFAR):", len(train_df))
print(" Novel-super count (superclass_index == NOVEL_SUPER_IDX):",
      (train_df["superclass_index"] == NOVEL_SUPER_IDX).sum())

# --- Build subclass -> superclass mapping from *combined* train_df ---
# This still satisfies "each subclass has a single superclass":
# - Original subclasses 0..86
# - Novel subclass 87 always maps to NOVEL_SUPER_IDX
sub_to_super = (
    train_df.groupby("subclass_index")["superclass_index"]
        .agg(lambda x: x.value_counts().index[0])
        .to_dict()
)

print("Example sub_to_super mapping (first 10):",
      dict(list(sub_to_super.items())[:10]))

```

Num superclasses: 4

```

Num subclasses: 88
CIFAR_NOVEL_MODE='large' → using 5000 CIFAR novel-super samples.
Combined train_df size (original + CIFAR): 11288
Novel-super count (superclass_index == NOVEL_SUPER_IDX): 5000
Example sub_to_super mapping (first 10): {0: 1, 1: 2, 2: 1, 3: 2, 4: 0, 5: 0, 6: 0, 7: 1, 8: 0, 9: 1}

```

```

[ ]: # Dataset functions
class BirdDogReptileDataset(Dataset):
    def __init__(self, df, img_dir, transform=None):
        self.df = df.reset_index(drop=True)
        self.img_dir = img_dir
        self.transform = transform

    def __len__(self):
        return len(self.df)

    def __getitem__(self, idx):
        row = self.df.iloc[idx]
        img_name = row["image"]
        img_path = os.path.join(self.img_dir, img_name)

        image = Image.open(img_path).convert("RGB")
        if self.transform:
            image = self.transform(image)

        super_idx = int(row["superclass_index"])
        sub_idx = int(row["subclass_index"])
        return image, super_idx, sub_idx

```

```

[ ]: # Test dataset (for leaderboard predictions)

class BirdDogReptileTestDataset(Dataset):
    def __init__(self, img_dir, transform=None):
        self.img_dir = img_dir
        self.transform = transform
        # assumes images are named 0.jpg, 1.jpg, ..., N-1.jpg
        self.filenames = sorted(os.listdir(img_dir), key=lambda x: int(os.path.splitext(x)[0]))

    def __len__(self):
        return len(self.filenames)

    def __getitem__(self, idx):
        img_name = self.filenames[idx]
        img_path = os.path.join(self.img_dir, img_name)

```

```

        image = Image.open(img_path).convert("RGB")
        if self.transform:
            image = self.transform(image)

    return image, img_name

```

[]: # Transforms

```

if DATA_AUGMENT:
    train_transform = transforms.Compose([
        transforms.Resize((IMG_SIZE, IMG_SIZE)),
        transforms.RandomHorizontalFlip(),
        transforms.RandomRotation(10),
        transforms.ToTensor(),
        transforms.Normalize(mean=[0.485, 0.456, 0.406],
                            std=[0.229, 0.224, 0.225]),
    ])
else:
    train_transform = transforms.Compose([
        transforms.Resize((IMG_SIZE, IMG_SIZE)),
        transforms.ToTensor(),
        transforms.Normalize(mean=[0.485, 0.456, 0.406],
                            std=[0.229, 0.224, 0.225]),
    ])

val_test_transform = transforms.Compose([
    transforms.Resize((IMG_SIZE, IMG_SIZE)),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406],
                        std=[0.229, 0.224, 0.225]),
])

```

[]: # Train/val split + loaders with optional pseudo-novel validation,
ensuring novel-super examples (super == NOVEL_SUPER_IDX) appear in both trainand val.

```

from math import ceil

pseudo_novel_loader = None      # default; will be set if USE_PSEUDO_NOVEL
heldout_subclasses = None       # to inspect later if needed

def split_seen_vs_novel_super(df, val_split, novel_super_idx, rng_seed=42):
    """
    Split dataframe into train/val, ensuring that both seen-super (0/1/2)
    and novel-super (== novel_super_idx) appear in both splits if present.
    """
    rng = np.random.default_rng(rng_seed)

```

```

df_novel_super = df[df["superclass_index"] == novel_super_idx]
df_seen_super = df[df["superclass_index"] != novel_super_idx]

print(" Seen-super samples:", len(df_seen_super))
print(" Novel-super samples:", len(df_novel_super))

# Split seen-super part
if len(df_seen_super) > 0:
    val_seen_size = int(len(df_seen_super) * val_split)
    val_seen_indices = rng.choice(len(df_seen_super), size=val_seen_size, u
→replace=False)
    val_seen_df = df_seen_super.iloc[val_seen_indices]
    train_seen_df = df_seen_super.drop(val_seen_df.index)
else:
    val_seen_df = df_seen_super.iloc[0:0]
    train_seen_df = df_seen_super.iloc[0:0]

# Split novel-super part (if any)
if len(df_novel_super) > 0:
    val_novel_size = max(1, int(len(df_novel_super) * val_split))
    val_novel_indices = rng.choice(len(df_novel_super), u
→size=val_novel_size, replace=False)
    val_novel_df = df_novel_super.iloc[val_novel_indices]
    train_novel_df = df_novel_super.drop(val_novel_df.index)
else:
    val_novel_df = df_novel_super.iloc[0:0]
    train_novel_df = df_novel_super.iloc[0:0]

# Combine splits and shuffle
train_split_df = pd.concat([train_seen_df, train_novel_df], u
→ignore_index=True)
val_split_df = pd.concat([val_seen_df, val_novel_df], u
→ignore_index=True)

train_split_df = train_split_df.sample(frac=1.0, random_state=rng_seed).
→reset_index(drop=True)
val_split_df = val_split_df.sample(frac=1.0, random_state=rng_seed + 1).
→reset_index(drop=True)

print(" Final split sizes:")
print("   train:", len(train_split_df))
print("   val: ", len(val_split_df))
print("   train novel-super:",
      (train_split_df["superclass_index"] == novel_super_idx).sum())
print("   val novel-super: ", "

```

```

        (val_split_df["superclass_index"] == novel_super_idx).sum())

    return train_split_df, val_split_df

if not USE_PSEUDO_NOVEL:
    # -----
    # Simple split, but novel-super-aware
    # -----
    print("[Simple split, novel-super aware]")

train_split_df, val_split_df = split_seen_vs_novel_super(
    train_df,
    VAL_SPLIT,
    NOVEL_SUPER_IDX,
    rng_seed=42,
)

train_dataset = BirdDogReptileDataset(
    train_split_df,
    TRAIN_IMG_DIR,
    transform=train_transform,
)
val_dataset = BirdDogReptileDataset(
    val_split_df,
    TRAIN_IMG_DIR,
    transform=val_test_transform,
)

train_loader = DataLoader(
    train_dataset,
    batch_size=BATCH_SIZE,
    shuffle=True,
    num_workers=NUM_WORKERS,
    pin_memory=True,
)
val_loader = DataLoader(
    val_dataset,
    batch_size=BATCH_SIZE,
    shuffle=False,
    num_workers=NUM_WORKERS,
    pin_memory=True,
)

print(f"[Simple split] Train size: {len(train_dataset)}, Val size:{len(val_dataset)}")

```

```

else:
    # -----
    # Pseudo-novel setup: hold out some subclasses entirely for pseudo-novel validation
    # while still keeping novel-super in both train and val.
    # -----


    print("[Pseudo-novel subclass split + novel-super-aware train/val]")

    # 1) choose subset of subclasses to treat as pseudo-novel
    all_subclasses = sorted(train_df["subclass_index"].unique())
    # Do NOT hold out the novel subclass itself (87)
    all_subclasses_no_novel = [c for c in all_subclasses if c != NOVEL_SUB_IDX]

    rng = np.random.default_rng(PSEUDO_NOVEL_SEED)

    num_holdout = max(1, int(len(all_subclasses_no_novel) * PSEUDO_NOVEL_FRACTION))
    heldout_subclasses = set(
        rng.choice(all_subclasses_no_novel, size=num_holdout, replace=False).
    tolist()
    )
    seen_subclasses = [c for c in all_subclasses if c not in heldout_subclasses]

    print(f"[Pseudo-novel] Total subclasses (excl. novel): {len(all_subclasses_no_novel)}")
    print(f"[Pseudo-novel] Held-out subclasses (pseudo-novel): {sorted(heldout_subclasses)}")
    print(f"[Pseudo-novel] Seen subclasses (incl. novel-sub): {len(seen_subclasses)}")

    # 2) split dataframe into seen vs pseudo-novel (by subclass)
    seen_mask = ~train_df["subclass_index"].isin(heldout_subclasses)
    seen_df = train_df[seen_mask].reset_index(drop=True)
    pseudo_novel_df = train_df[~seen_mask].reset_index(drop=True)

    print(f"[Pseudo-novel] Seen samples: {len(seen_df)}, Pseudo-novel samples: {len(pseudo_novel_df)}")

    # 3) Train/val split on seen data, but ensure novel-super in both
    train_split_df, val_split_df = split_seen_vs_novel_super(
        seen_df,
        VAL_SPLIT,
        NOVEL_SUPER_IDX,
        rng_seed=PSEUDO_NOVEL_SEED,
    )

```

```

train_dataset = BirdDogReptileDataset(
    train_split_df,
    TRAIN_IMG_DIR,
    transform=train_transform,
)
val_seen_dataset = BirdDogReptileDataset(
    val_split_df,
    TRAIN_IMG_DIR,
    transform=val_test_transform,
)

train_loader = DataLoader(
    train_dataset,
    batch_size=BATCH_SIZE,
    shuffle=True,
    num_workers=NUM_WORKERS,
    pin_memory=True,
)
val_loader = DataLoader(
    val_seen_dataset,
    batch_size=BATCH_SIZE,
    shuffle=False,
    num_workers=NUM_WORKERS,
    pin_memory=True,
)

# 4) pseudo-novel validation loader (all held-out subclasses, val-style
# transform)
pseudo_novel_dataset = BirdDogReptileDataset(
    pseudo_novel_df,
    TRAIN_IMG_DIR,
    transform=val_test_transform,
)
pseudo_novel_loader = DataLoader(
    pseudo_novel_dataset,
    batch_size=BATCH_SIZE,
    shuffle=False,
    num_workers=NUM_WORKERS,
    pin_memory=True,
)

print(f"[Pseudo-novel] Train size: {len(train_dataset)}, "
      f"Seen-val size: {len(val_seen_dataset)}, "
      f"Pseudo-novel val size: {len(pseudo_novel_dataset)}")

# Test loader is the same regardless

```

```

test_dataset = BirdDogReptileTestDataset(TEST_IMG_DIR,
                                         transform=val_test_transform)
test_loader = DataLoader(test_dataset, batch_size=64,
                        shuffle=False, num_workers=NUM_WORKERS)
# change batch size back to 1 if see any errors

print(f"Test size: {len(test_dataset)}")

[Pseudo-novel subclass split + novel-super-aware train/val]
[Pseudo-novel] Total subclasses (excl. novel): 87
[Pseudo-novel] Held-out subclasses (pseudo-novel): [1, 4, 14, 15, 17, 20, 27,
29, 39, 45, 51, 69, 71]
[Pseudo-novel] Seen subclasses (incl. novel-sub): 75
[Pseudo-novel] Seen samples: 10340, Pseudo-novel samples: 948
    Seen-super samples: 5340
    Novel-super samples: 5000
    Final split sizes:
        train: 9306
        val: 1034
        train novel-super: 4500
        val novel-super: 500
[Pseudo-novel] Train size: 9306, Seen-val size: 1034, Pseudo-novel val size: 948
Test size: 11180

```

4 SECTION 4: Backbone & Model Definitions

```

[ ]: # Backbone builder

# Using ImageNet pretrained ResNet backbone, and chopping off FC head to
# Transfer Learn

def build_resnet_backbone():
    # Using torchvision ResNet-18 with ImageNet weights
    if BACKBONE == "resnet18":
        base = models.resnet18(weights=models.ResNet18_Weights.IMAGENET1K_V1)
    elif BACKBONE == "resnet50":
        base = models.resnet50(weights=models.ResNet50_Weights.IMAGENET1K_V1)
    else:
        raise ValueError(f"Unknown BACKBONE: {BACKBONE}")
    # Remove the final classification layer
    in_features = base.fc.in_features
    base.fc = nn.Identity()
    return base, in_features

[ ]: # Methods for Shared backbone + two heads + KL divergence

class SharedBackboneTwoHeads(nn.Module):

```

```

def __init__(self, num_super, num_sub):
    super().__init__()
    self.backbone, feat_dim = build_resnet_backbone()
    self.super_head = nn.Linear(feat_dim, num_super)
    self.sub_head = nn.Linear(feat_dim, num_sub)

def forward(self, x):
    feats = self.backbone(x)
    super_logits = self.super_head(feats)
    sub_logits = self.sub_head(feats)
    return super_logits, sub_logits

```

[]: # KL helper that maps subclass probs to superclass probs using sub_to_super ↵mapping

```

def sub_probs_to_super_probs(sub_probs, sub_to_super, num_super):
    """
    sub_probs: (B, num_sub), softmax over subclasses
    returns: (B, num_super), summed probs per super-class
    """
    B, num_sub = sub_probs.shape
    super_probs = torch.zeros(B, num_super, device=sub_probs.device)

    for sub_idx, super_idx in sub_to_super.items():
        super_probs[:, super_idx] += sub_probs[:, sub_idx]

    # For safety: re-normalize in case of any numeric drift
    super_probs = super_probs / (super_probs.sum(dim=1, keepdim=True) + 1e-8)
    return super_probs

```

[]: # Single-head model for 2-separate models approach
Each model is just a backbone + one linear head

```

# to instantiate, one has num_classes = num_super, other has num_classes = num_sub

class SingleHeadModel(nn.Module):
    def __init__(self, num_classes):
        super().__init__()
        self.backbone, feat_dim = build_resnet_backbone()
        self.head = nn.Linear(feat_dim, num_classes)

    def forward(self, x):
        feats = self.backbone(x)
        logits = self.head(feats)
        return logits

```

5 SECTION 5: Training & evaluation utilities

```
[ ]: # functions for training

# we pass a flag mode to indicate which model using either:
# "two_heads_kl", "single_head_super" or "single_head"sub"

def accuracy_from_logits(logits, targets):
    preds = logits.argmax(dim=1)
    correct = (preds == targets).sum().item()
    total = targets.size(0)
    return correct, total

def train_one_epoch(model, optimizer, loader, criterion, mode, □
    ↪sub_to_super=None,
    num_super=None, alpha_kl=0.1, temperature=1.0):
    model.train()
    running_loss = 0.0
    super_correct = sub_correct = 0
    super_total = sub_total = 0

    for images, super_labels, sub_labels in loader:
        images = images.to(device)
        super_labels = super_labels.to(device)
        sub_labels = sub_labels.to(device)

        optimizer.zero_grad()

        if mode == "two_heads_kl":
            super_logits, sub_logits = model(images)
            # CE losses
            loss_super = criterion(super_logits, super_labels)
            loss_sub = criterion(sub_logits, sub_labels)

            # KL term between super head and aggregated subclass head
            with torch.no_grad():
                # target: super_probs
                super_probs = F.softmax(super_logits / temperature, dim=1)
                sub_probs = F.softmax(sub_logits / temperature, dim=1)
                agg_super_probs = sub_probs_to_super_probs(sub_probs, sub_to_super, □
                ↪num_super)

                # KL(super // agg_super) = sum p * (log p - log q)
                # using KLDivLoss with log_softmax input and probs target:
                kl_loss = F.kl_div(
```

```

        input=torch.log(agg_super_probs + 1e-8),
        target=super_probs,
        reduction="batchmean"
    )

    loss = loss_super + loss_sub + alpha_kl * kl_loss

    sc, st = accuracy_from_logits(super_logits, super_labels)
    suc, sut = accuracy_from_logits(sub_logits, sub_labels)
    super_correct += sc
    super_total += st
    sub_correct += suc
    sub_total += sut

    elif mode in ("single_head_super", "single_head_sub"):
        logits = model(images)
        if mode == "single_head_super":
            loss = criterion(logits, super_labels)
            sc, st = accuracy_from_logits(logits, super_labels)
            super_correct += sc
            super_total += st
        else:
            loss = criterion(logits, sub_labels)
            suc, sut = accuracy_from_logits(logits, sub_labels)
            sub_correct += suc
            sub_total += sut
    else:
        raise ValueError(f"Unknown mode {mode}")

    loss.backward()
    optimizer.step()

    running_loss += loss.item()

    avg_loss = running_loss / len(loader)
    metrics = {"loss": avg_loss}
    if super_total > 0:
        metrics["acc_super"] = super_correct / super_total
    if sub_total > 0:
        metrics["acc_sub"] = sub_correct / sub_total

    return metrics

@torch.no_grad()
def eval_one_epoch(model, loader, criterion, mode, sub_to_super=None,
                   num_super=None, alpha_kl=0.1, temperature=1.0):

```

```

model.eval()
running_loss = 0.0
super_correct = sub_correct = 0
super_total = sub_total = 0

for images, super_labels, sub_labels in loader:
    images = images.to(device)
    super_labels = super_labels.to(device)
    sub_labels = sub_labels.to(device)

    if mode == "two_heads_kl":
        super_logits, sub_logits = model(images)
        loss_super = criterion(super_logits, super_labels)
        loss_sub = criterion(sub_logits, sub_labels)

        with torch.no_grad():
            super_probs = F.softmax(super_logits / temperature, dim=1)
            sub_probs = F.softmax(sub_logits / temperature, dim=1)
            agg_super_probs = sub_probs_to_super_probs(sub_probs, sub_to_super, num_super)

        kl_loss = F.kl_div(
            input=torch.log(agg_super_probs + 1e-8),
            target=super_probs,
            reduction="batchmean"
        )

    loss = loss_super + loss_sub + alpha_kl * kl_loss

    sc, st = accuracy_from_logits(super_logits, super_labels)
    suc, sut = accuracy_from_logits(sub_logits, sub_labels)
    super_correct += sc
    super_total += st
    sub_correct += suc
    sub_total += sut

elif mode in ("single_head_super", "single_head_sub"):
    logits = model(images)
    if mode == "single_head_super":
        loss = criterion(logits, super_labels)
        sc, st = accuracy_from_logits(logits, super_labels)
        super_correct += sc
        super_total += st
    else:
        loss = criterion(logits, sub_labels)
        suc, sut = accuracy_from_logits(logits, sub_labels)
        sub_correct += suc

```

```

        sub_total += sut

    running_loss += loss.item()

    avg_loss = running_loss / len(loader)
    metrics = {"val_loss": avg_loss}
    if super_total > 0:
        metrics["val_acc_super"] = super_correct / super_total
    if sub_total > 0:
        metrics["val_acc_sub"] = sub_correct / sub_total

    return metrics

```

```

[ ]: # Analysis and Visualize: for Novel Subclass fine tuning
# Determines optimal threshold value from training data (except held out) vs. ↴
    ↴held out (pseudo-novel) subclass data

@torch.no_grad()
def collect_max_probs_sub(model, loader, mode="two_heads"):
    """
    Collect max softmax probabilities from the subclass head.

    mode:
        - "two_heads": model(images) -> (super_logits, sub_logits)
        - "sub_single_head": model(images) -> sub_logits
    """
    model.eval()
    probs = []

    for batch in loader:
        images = batch[0].to(device)

        if mode == "two_heads":
            _, sub_logits = model(images)
        elif mode == "sub_single_head":
            sub_logits = model(images)
        else:
            raise ValueError(f"Unknown mode: {mode}")

        # Softmax over subclasses, then take max per sample
        p = F.softmax(sub_logits, dim=1).max(dim=1).values
        probs.extend(p.cpu().numpy().tolist())

    return np.array(probs)

@torch.no_grad()
def analyze_tau_sub(model, mode="two_heads"):

```

```

"""
Compare max subclass probabilities on:
- seen validation (val_loader)
- pseudo-novel validation (pseudo_novel_loader)
and suggest TAU_SUB candidates.
"""

if pseudo_novel_loader is None:
    print("pseudo_novel_loader is None. "
          "Set USE_PSEUDO_NOVEL = True before building loaders to use this")
    return

# Collect max probs
seen_probs = collect_max_probs_sub(model, val_loader, mode=mode)
pseudo_probs = collect_max_probs_sub(model, pseudo_novel_loader, mode=mode)

# Summary stats
def summarize(name, arr):
    print(f"\n{name} subclass max-prob stats:")
    print(f"  count = {len(arr)}")
    print(f"  mean  = {arr.mean():.3f}")
    print(f"  std   = {arr.std():.3f}")
    percentiles = [1, 5, 10, 25, 50, 75, 90, 95, 99]
    pvals = {p: float(np.percentile(arr, p)) for p in percentiles}
    print("  percentiles:")
    for p in percentiles:
        print(f"    p{p:>2}: {pvals[p]:.3f}")
    return pvals

seen_p = summarize("Seen", seen_probs)
pseudo_p = summarize("Pseudo-novel", pseudo_probs)

# Simple candidate thresholds
# 1) 10th percentile of seen (reject only the lowest-confidence seen examples)
tau_candidate_1 = seen_p[10]

# 2) Midpoint between mean(seen) and mean(pseudo)
tau_candidate_2 = 0.5 * (seen_probs.mean() + pseudo_probs.mean())

print("\nSuggested TAU_SUB candidates:")
print(f"  tau_sub 10th percentile of seen: {tau_candidate_1:.3f}")
print(f"  tau_sub mean(seen + pseudo)/2: {tau_candidate_2:.3f}")
print("\nYou can start with one of these for TAU_SUB and adjust based on"
      "leaderboard/behavior.")

# Histogram visualization

```

```

plt.figure(figsize=(8, 5))
plt.hist(seen_probs, bins=30, alpha=0.5, label="Seen subclasses")
plt.hist(pseudo_probs, bins=30, alpha=0.5, label="Pseudo-novel subclasses")
plt.axvline(tau_candidate_1, linestyle="--", label=f"p10 seen ~\u2192{tau_candidate_1:.2f}")
plt.axvline(tau_candidate_2, linestyle=":", label=f"mean midpoint ~\u2192{tau_candidate_2:.2f}")
plt.xlabel("Max softmax probability (subclass head)")
plt.ylabel("Count")
plt.title("Subclass max-prob: seen vs pseudo-novel")
plt.legend()
plt.tight_layout()
plt.show()

# Optional: log histograms and candidates to W&B
if USE_WANDB:
    try:
        import wandb
        wandb.log({
            "sub_seen_maxprob_hist": wandb.Histogram(seen_probs),
            "sub_pseudo_maxprob_hist": wandb.Histogram(pseudo_probs),
            "tau_sub_candidate_p10_seen": tau_candidate_1,
            "tau_sub_candidate_mean_midpoint": tau_candidate_2,
        })
        print("Logged histograms and tau_sub candidates to Weights & Biases.")
    except Exception as e:
        print("Could not log to W&B:", e)

```

```

[ ]: # Analysis and Visualize: for Novel Superclass fine tuning
# Determines the optimal threshold from provided training data vs. Novel Super
    ↵data (more images)

@torch.no_grad()
def collect_max_probs_super(model, loader, mode="two_heads"):
    """
    Collect max softmax probabilities from the superclass head.

    mode:
        - "two_heads": model(images) -> (super_logits, sub_logits)
        - "super_single_head": model(images) -> super_logits
    """
    model.eval()
    probs = []

    for batch in loader:
        images = batch[0].to(device)

```

```

    if mode == "two_heads":
        super_logits, _ = model(images)
    elif mode == "super_single_head":
        super_logits = model(images)
    else:
        raise ValueError(f"Unknown mode: {mode}")

    p = F.softmax(super_logits, dim=1).max(dim=1).values
    probs.extend(p.cpu().numpy().tolist())

return np.array(probs)

@torch.no_grad()
def analyze_tau_super(model, mode="two_heads"):
    """
    Analyze max superclass probabilities on validation set, split into:
    - seen superclasses (super != NOVEL_SUPER_IDX)
    - novel superclasses (super == NOVEL_SUPER_IDX)

    This assumes:
    - val_loader batches look like (images, super_labels, sub_labels, ...)
    - NOVEL_SUPER_IDX is defined (e.g. 3)
    """
    model.eval()

    seen_probs = []
    novel_probs = []

    for batch in val_loader:
        images = batch[0].to(device)
        super_labels = batch[1].to(device)    # assumes (images, super, sub, ...)

        # Forward
        if mode == "two_heads":
            super_logits, _ = model(images)
        elif mode == "super_single_head":
            super_logits = model(images)
        else:
            raise ValueError(f"Unknown mode: {mode}")

        probs = F.softmax(super_logits, dim=1)          # (B, num_super)
        max_probs, _ = probs.max(dim=1)                 # (B,)

        novel_mask = (super_labels == NOVEL_SUPER_IDX)
        seen_mask = ~novel_mask

```

```

    if seen_mask.any():
        seen_probs.extend(max_probs[seen_mask].detach().cpu().numpy().
                           tolist())
    if novel_mask.any():
        novel_probs.extend(max_probs[novel_mask].detach().cpu().numpy().
                           tolist())

seen_probs = np.array(seen_probs)
novel_probs = np.array(novel_probs)

def summarize(name, arr):
    print(f"\n{name} superclass max-prob stats:")
    print(f"  count = {len(arr)}")
    print(f"  mean   = {arr.mean():.3f}")
    print(f"  std    = {arr.std():.3f}")
    percentiles = [1, 5, 10, 25, 50, 75, 90, 95, 99]
    pvals = {p: float(np.percentile(arr, p)) for p in percentiles}
    print("  percentiles:")
    for p in percentiles:
        print(f"    p{p:>2}: {pvals[p]:.3f}")
    return pvals

if len(seen_probs) == 0:
    print("No seen-super examples found in val_loader (super != NOVEL_SUPER_IDX).")
    return
seen_p = summarize("Seen superclasses (0/1/2)", seen_probs)

if len(novel_probs) == 0:
    print("\nNo novel-super examples (super == NOVEL_SUPER_IDX) in val_loader yet.")
    novel_p = None
else:
    novel_p = summarize("Novel superclasses (== NOVEL_SUPER_IDX)", novel_probs)

# ---- Candidate thresholds ----
# Start from seen distribution:
tau_p10_seen = seen_p[10]
tau_p5_seen = seen_p[5]
tau_mean_minus_std = max(0.0, min(1.0, seen_probs.mean() - seen_probs.
                                   std()))

# If we have novel-super stats, try to place tau between seen & novel means/
# percentiles
if novel_p is not None:

```

```

# Midpoint between seen mean and novel mean
tau_mean_mid = 0.5 * (seen_probs.mean() + novel_probs.mean())
print("\nSuggested TAU_SUPER candidates (using seen + novel):")
print(f" tau_super 10th percentile of seen:      {tau_p10_seen:.3f}")
print(f" tau_super 5th percentile of seen:       {tau_p5_seen:.3f}")
print(f" tau_super mean(seen) - std(seen):     □")
print(f" tau_super mean(seen & novel) midpoint: {tau_mean_mid:.3f}")
tau_candidates = [tau_p10_seen, tau_p5_seen, tau_mean_minus_std, □
tau_mean_mid]
else:
    print("\nSuggested TAU_SUPER candidates (seen only):")
    print(f" tau_super 10th percentile of seen:      {tau_p10_seen:.3f}")
    print(f" tau_super 5th percentile of seen:       {tau_p5_seen:.3f}")
    print(f" tau_super mean(seen) - std(seen):     □")
    print(f" tau_mean_minus_std: {tau_mean_minus_std:.3f}")
    tau_candidates = [tau_p10_seen, tau_p5_seen, tau_mean_minus_std]

# ---- Histogram plot ----
plt.figure(figsize=(8, 5))
plt.hist(seen_probs, bins=30, alpha=0.6, label="Seen superclasses (0/1/2)")
if len(novel_probs) > 0:
    plt.hist(novel_probs, bins=30, alpha=0.6, label="Novel superclasses")
    (3))

# Draw candidate lines (use a couple of them for visual reference)
for tau in tau_candidates[:3]:
    plt.axvline(tau, linestyle="--", alpha=0.7)

plt.xlabel("Max softmax probability (superclass head)")
plt.ylabel("Count")
plt.title("Superclass max-prob: seen vs novel-super on validation")
plt.legend()
plt.tight_layout()
plt.show()

# ---- Optional: log to W&B ----
if USE_WANDB:
    try:
        import wandb
        log_dict = {
            "super_seen_maxprob_hist": wandb.Histogram(seen_probs),
        }
    
```

```

    if len(novel_probs) > 0:
        log_dict["super_novel_maxprob_hist"] = wandb.
        ↪Histogram(novel_probs)
        # Also log first few tau candidates
        log_dict["tau_super_candidate_1"] = float(tau_candidates[0])
        if len(tau_candidates) > 1:
            log_dict["tau_super_candidate_2"] = float(tau_candidates[1])
        if len(tau_candidates) > 2:
            log_dict["tau_super_candidate_3"] = float(tau_candidates[2])
        wandb.log(log_dict)
        print("Logged superclass histograms and tau_super candidates to"
        ↪Weights & Biases.)
    except Exception as e:
        print("Could not log to W&B:", e)

```

[]: # Additional eval helper functions

```

# This calculates the following:
# 1. How often we correctly keep seen classes (low false "novel" rate)
# 2. How often we correctly identify CIFAR "novel super" samples (new data)

```

```

@torch.no_grad()
def evaluate_on_val_with_novelty(model, mode="two_heads",
                                  tau_super=TAU_SUPER, tau_sub=TAU_SUB,
                                  loader=None, name="val"):

"""

```

Evaluate a model on a loader with novel thresholds applied.

For subclasses, we ONLY care about:

- overall subclass acc
- seen-subclass acc / false-novel rate

We DO NOT report "novel-subclass accuracy" here because there are no ground-truth novel-sub labels in the original validation data.

"""

```

if loader is None:
    loader = val_loader

```

```

model.eval()

```

```

total = 0
super_correct = 0
sub_correct = 0

```

```

seen_super_correct = seen_super_total = 0
novel_super_correct = novel_super_total = 0

```

```

seen_sub_correct = seen_sub_total = 0
novel_sub_correct = novel_sub_total = 0 # still counted but not reported

for batch in loader:
    images = batch[0].to(device)
    super_true = batch[1].to(device)
    sub_true = batch[2].to(device)

    if mode == "two_heads":
        super_logits, sub_logits = model(images)
    elif mode == "super_single_head":
        super_logits = model(images); sub_logits = None
    elif mode == "sub_single_head":
        super_logits = None; sub_logits = model(images)
    else:
        raise ValueError(f"Unknown mode: {mode}")

batch_size = images.size(0)
total += batch_size

# --- Superclass ---
if super_logits is not None:
    super_probs = F.softmax(super_logits, dim=1)
    max_super_probs, super_idx = super_probs.max(dim=1)

    super_pred = super_idx.clone()
    novel_mask_super = max_super_probs < tau_super
    super_pred[novel_mask_super] = NOVEL_SUPER_IDX

    match_super = (super_pred == super_true)
    super_correct += match_super.sum().item()

    seen_mask_super = (super_true != NOVEL_SUPER_IDX)
    novel_mask_super_label = (super_true == NOVEL_SUPER_IDX)

    if seen_mask_super.any():
        seen_super_correct += match_super[seen_mask_super].sum().item()
        seen_super_total += seen_mask_super.sum().item()
    if novel_mask_super_label.any():
        novel_super_correct += match_super[novel_mask_super_label].
            sum().item()
        novel_super_total += novel_mask_super_label.sum().item()

# --- Subclass ---
if sub_logits is not None:
    sub_probs = F.softmax(sub_logits, dim=1)
    max_sub_probs, sub_idx = sub_probs.max(dim=1)

```

```

    sub_pred = sub_idx.clone()
    novel_mask_sub = max_sub_probs < tau_sub
    sub_pred[novel_mask_sub] = NOVEL_SUB_IDX

    match_sub = (sub_pred == sub_true)
    sub_correct += match_sub.sum().item()

    seen_mask_sub = (sub_true != NOVEL_SUB_IDX)
    novel_mask_sub_label = (sub_true == NOVEL_SUB_IDX)

    if seen_mask_sub.any():
        seen_sub_correct += match_sub[seen_mask_sub].sum().item()
        seen_sub_total += seen_mask_sub.sum().item()
    if novel_mask_sub_label.any():
        # counted but not reported; there shouldn't be any in original
        ↵val
        novel_sub_correct += match_sub[novel_mask_sub_label].sum().
    ↵item()
        novel_sub_total += novel_mask_sub_label.sum().item()

# ----- Compute metrics dict -----
metrics = {}

if total > 0:
    if super_correct > 0 or seen_super_total + novel_super_total > 0:
        metrics["overall_super_acc"] = super_correct / total
    if sub_correct > 0 or seen_sub_total + novel_sub_total > 0:
        metrics["overall_sub_acc"] = sub_correct / total

if seen_super_total > 0:
    metrics["seen_super_acc"] = seen_super_correct / seen_super_total
    metrics["seen_super_false_novel"] = 1.0 - metrics["seen_super_acc"]
if novel_super_total > 0:
    metrics["novel_super_acc"] = novel_super_correct / novel_super_total

if seen_sub_total > 0:
    metrics["seen_sub_acc"] = seen_sub_correct / seen_sub_total
    metrics["seen_sub_false_novel"] = 1.0 - metrics["seen_sub_acc"]
# NOTE: we deliberately do NOT add novel-sub metrics to `metrics`,
# because they are not meaningful for this dataset.

# ----- Print summary -----
print(f"\n==== Evaluation on {name} ===")
if "overall_super_acc" in metrics:
    print(f"Overall superclass acc: {metrics['overall_super_acc']:.4f}")
if "overall_sub_acc" in metrics:

```

```

        print(f"Overall subclass acc: {metrics['overall_sub_acc']:.4f}")

    if "seen_super_acc" in metrics:
        print(f"Seen superclass acc (true super != novel): {metrics['seen_super_acc']:.4f}")
        print(f"Seen superclass false-novel rate: {metrics['seen_super_false_novel']:.4f}")

    if "novel_super_acc" in metrics:
        print(f"Novel superclass acc (true super == novel): {metrics['novel_super_acc']:.4f}")

    if "seen_sub_acc" in metrics:
        print(f"Seen subclass acc (true sub != novel): {metrics['seen_sub_acc']:.4f}")
        print(f"Seen subclass false-novel rate: {metrics['seen_sub_false_novel']:.4f}")

# ----- Optional: log to W&B (only real metrics) -----
if USE_WANDB:
    log_dict = {}
    for k, v in metrics.items():
        log_dict[f"{name}_{k}"] = v
    if log_dict:
        wandb.log(log_dict)

return metrics

# This enables us to evaluate how often we correctly mark held-out subclasses as novel (proxy for leaderboard performance on subclasses)
@torch.no_grad()
def evaluate_pseudo_novel_sub_with_novelty(model, mode="two_heads",
                                             tau_sub=TAU_SUB,
                                             loader=None,
                                             name="pseudo_novel_sub"):
    """
    Evaluate how well the model flags held-out subclasses as novel.

    Assumes:
    - loader yields only *held-out* subclasses (true unseen subclasses)
    - true labels are NOT NOVEL_SUB_IDX, but we *want* the model to predict NOVEL_SUB_IDX.
    """

    if loader is None:

```

```

        loader = pseudo_novel_loader

    if loader is None:
        print("No pseudo_novel_loader available.")
        return {}

model.eval()

total = 0
predicted_novel = 0
predicted_seen = 0

for batch in loader:
    images = batch[0].to(device)
    # we don't actually need the labels here for correctness, only for
    ↪counting

    if mode == "two_heads":
        _, sub_logits = model(images)
    elif mode == "sub_single_head":
        sub_logits = model(images)
    else:
        raise ValueError(f"Unknown mode for pseudo-novel eval: {mode}")

    sub_probs = F.softmax(sub_logits, dim=1)
    max_sub_probs, sub_idx = sub_probs.max(dim=1)

    sub_pred = sub_idx.clone()
    novel_mask_sub = max_sub_probs < tau_sub
    sub_pred[novel_mask_sub] = NOVEL_SUB_IDX

    batch_size = images.size(0)
    total += batch_size

    predicted_novel += (sub_pred == NOVEL_SUB_IDX).sum().item()
    predicted_seen += (sub_pred != NOVEL_SUB_IDX).sum().item()

metrics = []
if total > 0:
    metrics["pseudo_novel_sub_novel_rate"] = predicted_novel / total
    metrics["pseudo_novel_sub_false_seen"] = predicted_seen / total

print(f"\n==== Evaluation on {name} (held-out subclasses) ===")
if total > 0:
    print(f"Fraction flagged as novel (good):"
    ↪{metrics['pseudo_novel_sub_novel_rate']:.4f}")

```

```

        print(f"Fraction mapped to seen subclasses (bad): {metrics['pseudo_novel_sub_false_seen']:.4f}")

    if USE_WANDB and metrics:
        log_dict = {f"{name}_{k}": v for k, v in metrics.items()}
        wandb.log(log_dict)

    return metrics

```

[]: # Dashboard for our evals so we can determine how model will likely perform on leaderboard evaluation

```

@torch.no_grad()
def novelty_dashboard(model, mode="two_heads",
                      tau_super=TAU_SUPER, tau_sub=TAU_SUB,
                      include_pseudo=True):
    """
    Run thresholded evals and show a compact table of key metrics.

    Subclass side focuses on:
    - seen subclasses (false-novel rate)
    - held-out pseudo-novel subclasses (how often marked as novel)
    """
    rows = []

    # --- Main val_loader stats (seen + CIFAR-novel) ---
    val_metrics = evaluate_on_val_with_novelty(
        model,
        mode=mode,
        tau_super=tau_super,
        tau_sub=tau_sub,
        loader=val_loader,
        name="val",
    )

    # Superclass (val)
    if "seen_super_acc" in val_metrics:
        rows.append({
            "Split": "val",
            "Head": "super",
            "Metric": "Seen superclass accuracy",
            "Meaning": "Correctly keep seen superclasses as seen",
            "Value": float(val_metrics["seen_super_acc"]),
        })
    if "seen_super_false_novel" in val_metrics:
        rows.append({
            "Split": "val",

```

```

        "Head": "super",
        "Metric": "Seen superclass false-novel rate",
        "Meaning": "Seen superclasses incorrectly flipped to novel",
        "Value": float(val_metrics["seen_super_false_novel"]),
    })

if "novel_super_acc" in val_metrics:
    rows.append({
        "Split": "val",
        "Head": "super",
        "Metric": "Novel superclass accuracy (CIFAR)",
        "Meaning": "CIFAR novel-super samples correctly predicted as novel",
        "Value": float(val_metrics["novel_super_acc"]),
    })

# Subclass (val) - ONLY seen-subclass metrics
if "seen_sub_acc" in val_metrics:
    rows.append({
        "Split": "val",
        "Head": "sub",
        "Metric": "Seen subclass accuracy",
        "Meaning": "Correctly keep seen subclasses as seen",
        "Value": float(val_metrics["seen_sub_acc"]),
    })

if "seen_sub_false_novel" in val_metrics:
    rows.append({
        "Split": "val",
        "Head": "sub",
        "Metric": "Seen subclass false-novel rate",
        "Meaning": "Seen subclasses incorrectly flipped to novel",
        "Value": float(val_metrics["seen_sub_false_novel"]),
    })

# NOTE: we deliberately do NOT add a "novel subclass accuracy" row here.

# --- Pseudo-novel subclass stats (held-out subclasses) ---
pseudo_metrics = []
if include_pseudo and pseudo_novel_loader is not None and mode in [
    ("two_heads", "sub_single_head")]:
    pseudo_metrics = evaluate_pseudo_novel_sub_with_novelty(
        model,
        mode=mode,
        tau_sub=tau_sub,
        loader=pseudo_novel_loader,
        name="pseudo_novel_sub",
    )

    if "pseudo_novel_sub_novel_rate" in pseudo_metrics:

```

```

rows.append({
    "Split": "pseudo_novel",
    "Head": "sub",
    "Metric": "Pseudo-novel marked as novel",
    "Meaning": "Held-out subclasses correctly flagged as novel",
    "Value": float(pseudo_metrics["pseudo_novel_sub_novel_rate"]),
})
if "pseudo_novel_sub_false_seen" in pseudo_metrics:
    rows.append({
        "Split": "pseudo_novel",
        "Head": "sub",
        "Metric": "Pseudo-novel mapped to seen",
        "Meaning": "Held-out subclasses wrongly mapped to some seen",
        "Value": float(pseudo_metrics["pseudo_novel_sub_false_seen"]),
    })

# --- Config / settings summary (for *printing* only) ---
config_rows = [
    {
        "Split": "config",
        "Head": "-",
        "Metric": "BACKBONE",
        "Meaning": "Feature extractor (e.g. resnet18 / resnet50)",
        "Value": BACKBONE,
    },
    {
        "Split": "config",
        "Head": "-",
        "Metric": "TAU_SUPER",
        "Meaning": "Novelty threshold for superclass head",
        "Value": str(tau_super),
    },
    {
        "Split": "config",
        "Head": "-",
        "Metric": "TAU_SUB",
        "Meaning": "Novelty threshold for subclass head",
        "Value": str(tau_sub),
    },
    {
        "Split": "config",
        "Head": "-",
        "Metric": "CIFAR_NOVEL_MODE",
        "Meaning": "Extra novel-super CIFAR data mode",
        "Value": str(CIFAR_NOVEL_MODE),
    },
]

```

```

{
    "Split": "config",
    "Head": "-",
    "Metric": "FINE_TUNE_MODE",
    "Meaning": "Backbone training mode (full vs frozen)",
    "Value": str(FINE_TUNE_MODE),
},
{
    "Split": "config",
    "Head": "-",
    "Metric": "APPROACH",
    "Meaning": "Model architecture (two_heads vs two_models)",
    "Value": str(APPROACH),
},
{
    "Split": "config",
    "Head": "-",
    "Metric": "USE_PSEUDO_NOVEL",
    "Meaning": "Using held-out subclasses for pseudo-novel eval",
    "Value": str(bool(USE_PSEUDO_NOVEL)),
},
{
    "Split": "config",
    "Head": "-",
    "Metric": "DATA_AUGMENT",
    "Meaning": "Whether data augmentation is enabled for training",
    "Value": str(bool(DATA_AUGMENT)),
},
],
]

rows.extend(config_rows)

if not rows:
    print("No metrics to display in dashboard.")
    return None

dashboard_df = pd.DataFrame(rows)
dashboard_df = dashboard_df.sort_values(
    by=["Split", "Head", "Metric"]
).reset_index(drop=True)

print("\n==== Novelty Dashboard ====")
print(dashboard_df)

# --- Log to W&B: metrics only, numeric Value column ---
if USE_WANDB:
    try:

```

```

wandb_table_df = dashboard_df.copy()

def _to_str(v):
    if isinstance(v, float):
        return f"{v:.6f}"
    return str(v)

wandb_table_df["Split"] = wandb_table_df["Split"].astype(str)
wandb_table_df["Head"] = wandb_table_df["Head"].astype(str)
wandb_table_df["Metric"] = wandb_table_df["Metric"].astype(str)
wandb_table_df["Meaning"] = wandb_table_df["Meaning"].astype(str)
wandb_table_df["Value"] = wandb_table_df["Value"].apply(_to_str)

wandb.log({"novelty_dashboard": wandb.
    ↪Table(dataframe=wandb_table_df)})
except Exception as e:
    print("Could not log dashboard table to W&B:", e)

return dashboard_df

```

[]: # Helper to choose backbone (freeze or full) and choose optimizer parameter

```

def setup_backbone_training(model, fine_tune_mode="full", lr_full=1e-4, ↪
    ↪lr_head=1e-3):
    """
    Given a model with attributes:
    - model.backbone (all feature extractor layers)
    - head layers (e.g. super/sub heads) as other modules,
    freeze or unfreeze the backbone and return an optimizer.

    Returns:
        optimizer, effective_lr
    """
    if fine_tune_mode == "full":
        # Everything trainable
        for p in model.parameters():
            p.requires_grad = True
        trainable_params = model.parameters()
        lr = lr_full
    elif fine_tune_mode == "frozen":
        # Freeze backbone, train only heads
        for p in model.backbone.parameters():
            p.requires_grad = False
        # Only parameters that still require grad will be optimized
        trainable_params = [p for p in model.parameters() if p.requires_grad]
        lr = lr_head

```

```

        print(f"Freezing backbone; training {len(trainable_params)} parameter\u20d7
        tensors in heads only.")
    else:
        raise ValueError(f"Unknown FINE_TUNE_MODE: {fine_tune_mode}")

optimizer = optim.Adam(trainable_params, lr=lr, weight_decay=WEIGHT_DECAY)
return optimizer

```

6 SECTION 6: Approach A: Shared backbone, two heads + KL

[]: # Original idea

```

LR = 1e-4 # maybe two different learning rates
ALPHA_KL = 0.1
TEMPERATURE = 1.0

if APPROACH == "two_heads":
    model_two_heads = SharedBackboneTwoHeads(
        num_super=num_super,
        num_sub=num_sub
    ).to(device)

    criterion = nn.CrossEntropyLoss()
    optimizer = setup_backbone_training(
        model_two_heads,
        fine_tune_mode=FINE_TUNE_MODE,
        lr_full=LR,
        lr_head=LR_HEAD,
    )

    run_two_heads = None
    if USE_WANDB:
        run_two_heads = wandb.init(
            entity=WANDB_ENTITY,
            project=WANDB_PROJECT,
            name=make_run_name("two_heads_kl_resnet"),
            group="two_heads", # to group/filter all two_heads runs
            config={
                "approach": "two_heads_kl",
                "backbone": BACKBONE,
                "epochs": EPOCHS,
                "lr_full": LR,
                "lr_head": LR_HEAD,
                "fine_tune_mode": FINE_TUNE_MODE,
                "alpha_kl": ALPHA_KL,
                "temperature": TEMPERATURE,

```

```

        "img_size": IMG_SIZE,
    },
)

best_val_score = 0.0

print("Training shared-backbone two-heads model (KL-coupled):")
for epoch in range(1, EPOCHS + 1):
    train_metrics = train_one_epoch(
        model_two_heads,
        optimizer,
        train_loader,
        criterion,
        mode="two_heads_kl",
        sub_to_super=sub_to_super,
        num_super=num_super,
        alpha_kl=ALPHA_KL,
        temperature=TEMPERATURE,
    )

    val_metrics = eval_one_epoch(
        model_two_heads,
        val_loader,
        criterion,
        mode="two_heads_kl",
        sub_to_super=sub_to_super,
        num_super=num_super,
        alpha_kl=ALPHA_KL,
        temperature=TEMPERATURE,
    )

    val_acc_super = val_metrics.get("val_acc_super", 0.0)
    val_acc_sub = val_metrics.get("val_acc_sub", 0.0)
    val_loss = val_metrics["val_loss"]

    print(
        f"[Two-heads] Epoch {epoch}: "
        f"train_loss={train_metrics['loss']:.4f}, "
        f"val_loss={val_loss:.4f}, "
        f"val_acc_super={val_acc_super:.4f}, "
        f"val_acc_sub={val_acc_sub:.4f}"
    )

if USE_WANDB:
    # prefix metrics so they don't collide with two-model ones
    log_dict = {
        "epoch": epoch,

```

```

        "two_heads_train_loss": train_metrics["loss"],
        "two_heads_val_loss": val_loss,
        "two_heads_val_acc_super": val_acc_super,
        "two_heads_val_acc_sub": val_acc_sub,
    }
wandb.log(log_dict, step=epoch)

# simple combined score: average of super/sub val accuracy
val_score = 0.5 * val_acc_super + 0.5 * val_acc_sub
if val_score > best_val_score:
    best_val_score = val_score
    torch.save(
        model_two_heads.state_dict(),
        os.path.join(DATA_ROOT, "best_two_heads_kl.pth"),
    )
    print(" Saved new best two-heads model")

best_ckpt_path = os.path.join(DATA_ROOT, "best_two_heads_kl.pth")
model_two_heads.load_state_dict(torch.load(best_ckpt_path, map_location=device))

analyze_tau_sub(model_two_heads, mode="two_heads")
analyze_tau_super(model_two_heads, mode="two_heads")
evaluate_on_val_with_novelty(model_two_heads, mode="two_heads",
                              tau_super=TAU_SUPER, tau_sub=TAU_SUB,
                              loader=val_loader, name="val_two_heads")
evaluate_pseudo_novel_sub_with_novelty(
    model_two_heads, mode="two_heads", tau_sub=TAU_SUB
)
novelty_dashboard(model_two_heads, mode="two_heads",
                  tau_super=TAU_SUPER, tau_sub=TAU_SUB)

if run_two_heads is not None:
    run_two_heads.finish()

else:
    print("APPROACH is not 'two_heads'; skipping two-heads training in this cell.")

```

APPROACH is not 'two_heads'; skipping two-heads training in this cell.

7 Aksel Modifications

7.1 Newer Resnet

```
[ ]: def build_resnet_backbone(backbone: str=BACKBONE):
    # Using torchvision ResNet-18 with ImageNet weights
    if backbone == "resnet18":
        base = models.resnet18(weights=models.ResNet18_Weights.IMAGENET1K_V1)
    elif backbone == "resnet50":
        base = models.resnet50(weights=models.ResNet50_Weights.IMAGENET1K_V2)
    else:
        raise ValueError(f"Unknown BACKBONE: {backbone}")
    # Remove the final classification layer
    in_features = base.fc.in_features
    base.fc = nn.Identity() # type: ignore[assignment]
    return base, in_features
```

7.2 New training routine

```
[ ]: def train_one_epoch(
    model: nn.Module,
    optimizer: torch.optim.Optimizer,
    loader: torch.utils.data.DataLoader,
    criterion: nn.Module,
    mode: str,
    sub_to_super: dict|None = None,
    num_super: int | None = None,
    trained_superclass_model: nn.Module | None = None,
    alpha_kl=0.1,
    temperature=1.0,
):
    """
    Optionally, add super class to calculate KL divergence penalty
    """
    model.train()
    running_loss = 0.0
    super_correct = sub_correct = 0
    super_total = sub_total = 0
    loss: torch.Tensor

    for images, super_labels, sub_labels in loader:
        images: torch.Tensor
        super_labels: torch.Tensor
        sub_labels: torch.Tensor
        images = images.to(device)
        super_labels = super_labels.to(device)
        sub_labels = sub_labels.to(device)
```

```

optimizer.zero_grad()

if mode == "two_heads_kl":
    super_logits, sub_logits = model(images)
    # CE losses
    loss_super = criterion(super_logits, super_labels)
    loss_sub = criterion(sub_logits, sub_labels)

    # KL term between super head and aggregated subclass head
    with torch.no_grad():
        # target: super_probs
        super_probs = F.softmax(super_logits / temperature, dim=1)
        sub_probs = F.softmax(sub_logits / temperature, dim=1)
        agg_super_probs = sub_probs_to_super_probs(sub_probs, sub_to_super, ↵
                                                    num_super)

        #  $KL(super // agg\_super) = \sum p * (\log p - \log q)$ 
        # using KLDivLoss with log_softmax input and probs target:
        kl_loss = F.kl_div(
            F.log_softmax(super_logits / temperature, dim=1),
            (agg_super_probs + 1e-8).log(),
            reduction="batchmean",
            log_target=True,
        )

    loss = loss_super + loss_sub + alpha_kl * kl_loss

    sc, st = accuracy_from_logits(super_logits, super_labels)
    suc, sut = accuracy_from_logits(sub_logits, sub_labels)
    super_correct += sc
    super_total += st
    sub_correct += suc
    sub_total += sut

elif mode in ("single_head_super", "single_head_sub"):
    logits: torch.Tensor = model(images)
    if mode == "single_head_super":
        loss = criterion(logits, super_labels)
        sc, st = accuracy_from_logits(logits, super_labels)
        super_correct += sc
        super_total += st
    else:
        # Base subclass CE loss
        loss_sub = criterion(logits, sub_labels)

        # Optional: superclass-consistency penalty using a pretrained ↵
        # superclass "teacher"

```

```

# Idea: teacher gives  $p(\text{super}/x)$ . We aggregate the subclass ↵
↳ head's probabilities into
    #  $q(\text{super}/x)$  by summing subclass probs within each super. ↵
↳
    ↳ Penalize mismatch.
        loss_consistency = torch.zeros(() , device=device)

        if trained_superclass_model is not None:
            assert sub_to_super is not None, "sub_to_super mapping ↵
            ↳ required for superclass-consistency"

        with torch.no_grad():
            super_teacher_logits = trained_superclass_model(images)
            super_teacher_probs = F.softmax(super_teacher_logits, ↵
            ↳ dim=1)

            # Aggregate subclass probs to super probs
            sub_probs = F.softmax(logits, dim=1)
            agg_super_probs = sub_probs_to_super_probs(sub_probs, ↵
            ↳ sub_to_super, num_super)

            # Consistency loss option A (soft):  $KL(\text{teacher} // \text{agg\_super})$ 
            # Avoid 0 * log 0 by clamping.
            eps = 1e-8
            loss_consistency = F.kl_div(
                (agg_super_probs.clamp_min(eps)).log(),
                super_teacher_probs,
                reduction="batchmean",
            )

            # Consistency loss option B (hard): encourage probability ↵
            ↳ mass to fall inside the teacher's top super
                # Uncomment to add this as well.
                # in_super_prob = agg_super_probs.gather(1, ↵
            ↳ super_pred_indices.unsqueeze(1)).squeeze(1)
                # loss_consistency = loss_consistency + (-torch.
            ↳ log(in_super_prob.clamp_min(eps))).mean()

            # Weight the consistency term (tune as needed)
            loss = loss_sub + ALPHA_SUPER_CONSISTENCY * loss_consistency

            suc, sut = accuracy_from_logits(logits, sub_labels)
            sub_correct += suc
            sub_total += sut
        else:
            raise ValueError(f"Unknown mode {mode}")

```

```

        loss.backward()
        optimizer.step()

        running_loss += loss.item()

    avg_loss = running_loss / len(loader)
    metrics = {"loss": avg_loss}
    if super_total > 0:
        metrics["acc_super"] = super_correct / super_total
    if sub_total > 0:
        metrics["acc_sub"] = sub_correct / sub_total

    return metrics

```

7.3 New Single Head

```
[ ]: # Single-head model for 2-separate models approach
# Each model is just a backbone + one linear head

# to instantiate, one has num_classes = num_super, other has num_classes = num_sub

class CosineClassifier(nn.Module):
    """
    Cosine-similarity classifier (normalized linear head).
    logits = scale * cos(theta), where cos(theta)=<normalize(x), normalize(W)>.
    """

    def __init__(self, in_features: int, out_features: int, scale: float = 30.0, learn_scale: bool = True):
        super().__init__()
        self.weight = nn.Parameter(torch.empty(out_features, in_features))
        nn.init.xavier_uniform_(self.weight)

        if learn_scale:
            self.scale = nn.Parameter(torch.tensor(float(scale)))
        else:
            self.register_buffer("scale", torch.tensor(float(scale)))

    def forward(self, x: torch.Tensor) -> torch.Tensor:
        x = F.normalize(x, p=2, dim=1) # [B, D]
        w = F.normalize(self.weight, p=2, dim=1) # [C, D]
        logits = x @ w.t() # [B, C]
        return logits * self.scale

```

```

class SingleHeadModel(nn.Module):
    def __init__(
        self,
        num_classes: int,
        head: str, # "linear" or "cosine"
        cosine_scale: float = 30.0,
        learn_scale: bool = True,
    ):
        super().__init__()
        self.backbone, feat_dim = build_resnet_backbone()

        if head == "cosine":
            self.head = nn.Sequential(
                CosineClassifier(
                    in_features=feat_dim,
                    out_features=(out := num_classes),
                    scale=cosine_scale,
                    learn_scale=learn_scale,
                ),
                nn.ReLU(),
                nn.Linear(out, num_classes),
            )
        elif head == "linear":
            self.head = nn.Linear(feat_dim, num_classes)
    def forward(self, x: torch.Tensor) -> torch.Tensor:
        feats = self.backbone(x)
        logits = self.head(feats)
        return logits

```

8 Print Vars

```
[ ]: SUPER_HEAD_TYPE = "linear" # "linear" or "cosine"
SUB_HEAD_TYPE = "cosine" # "linear" or "cosine"
ALPHA_SUPER_CONSISTENCY = 0.2 # unused
USE_PSEUDO_NOVEL = False
```

```
[ ]: print(BACKBONE, FINE_TUNE_MODE, SUPER_HEAD_TYPE, SUB_HEAD_TYPE, USE_PSEUDO_NOVEL)
```

resnet50 full linear cosine True

9 SECTION 7: Approach B - Two separate models

9.1 Train Super

```
[ ]: # Second idea to compare

LR = 1e-4

if APPROACH == "two_models":
    # Superclass model
    model_super = SingleHeadModel(num_classes=num_super, head=SUPER_HEAD_TYPE).
    ↪to(device)
    criterion_super = nn.CrossEntropyLoss()
    optimizer_super = setup_backbone_training(
        model_super,
        fine_tune_mode=FINE_TUNE_MODE,
        lr_full=LR,
        lr_head=LR_HEAD,
    )

    run_super = None
    if USE_WANDB:
        run_super = wandb.init(
            entity=WANDB_ENTITY,
            project=WANDB_PROJECT,
            name=make_run_name("two_models_super_resnet"),
            group="two_models",    # to filter all two_models runs together
            config={
                "approach": "two_models",
                "head": "super",
                "backbone": BACKBONE,
                "epochs": EPOCHS,
                "lr_full": LR,
                "lr_head": LR_HEAD,
                "fine_tune_mode": FINE_TUNE_MODE,
                "img_size": IMG_SIZE,
                'super_head_type': SUPER_HEAD_TYPE,
            },
        )
    best_val_super = 0.0

    print("Training superclass model:")
    for epoch in range(1, EPOCHS + 1):
        train_metrics = train_one_epoch(
            model_super,
            optimizer_super,
```

```

        train_loader,
        criterion_super,
        mode="single_head_super",
    )
    val_metrics = eval_one_epoch(
        model_super,
        val_loader,
        criterion_super,
        mode="single_head_super",
    )

    val_acc = val_metrics.get("val_acc_super", 0.0)
    print(
        f"[Super] Epoch {epoch}: "
        f"train_loss={train_metrics['loss']:.4f}, "
        f"val_loss={val_metrics['val_loss']:.4f}, "
        f"val_acc_super={val_acc:.4f}"
    )

    if USE_WANDB:
        wandb.log(
            {
                "epoch": epoch,
                "super_train_loss": train_metrics["loss"],
                "super_val_loss": val_metrics["val_loss"],
                "super_val_acc": val_acc,
            },
            step=epoch,
        )

    if val_acc > best_val_super:
        best_val_super = val_acc
        torch.save(
            model_super.state_dict(),
            os.path.join(DATA_ROOT, "best_super_model.pth"),
        )
        print(" Saved new best superclass model")
    if val_acc == 1.0:
        print('ACHIEVED PERFECT SCORE.')

    if run_super is not None:
        run_super.finish()

<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>

```

```

<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>

Training superclass model:
[Super] Epoch 1: train_loss=0.2098, val_loss=0.0126, val_acc_super=0.9932
    Saved new best superclass model
[Super] Epoch 2: train_loss=0.0141, val_loss=0.0059, val_acc_super=0.9981
    Saved new best superclass model
[Super] Epoch 3: train_loss=0.0077, val_loss=0.0042, val_acc_super=0.9990
    Saved new best superclass model
[Super] Epoch 4: train_loss=0.0044, val_loss=0.0016, val_acc_super=1.0000
    Saved new best superclass model
ACHIEVED PERFECT SCORE.

[Super] Epoch 5: train_loss=0.0033, val_loss=0.0036, val_acc_super=0.9990
[Super] Epoch 6: train_loss=0.0037, val_loss=0.0045, val_acc_super=0.9990
[Super] Epoch 7: train_loss=0.0021, val_loss=0.0077, val_acc_super=0.9981
[Super] Epoch 8: train_loss=0.0016, val_loss=0.0018, val_acc_super=0.9990
[Super] Epoch 9: train_loss=0.0018, val_loss=0.0082, val_acc_super=0.9971
[Super] Epoch 10: train_loss=0.0007, val_loss=0.0012, val_acc_super=1.0000
ACHIEVED PERFECT SCORE.

[Super] Epoch 11: train_loss=0.0029, val_loss=0.0012, val_acc_super=1.0000
ACHIEVED PERFECT SCORE.

[Super] Epoch 12: train_loss=0.0026, val_loss=0.0008, val_acc_super=1.0000
ACHIEVED PERFECT SCORE.

[Super] Epoch 13: train_loss=0.0006, val_loss=0.0017, val_acc_super=0.9990
[Super] Epoch 14: train_loss=0.0059, val_loss=0.0008, val_acc_super=1.0000
ACHIEVED PERFECT SCORE.

[Super] Epoch 15: train_loss=0.0033, val_loss=0.0004, val_acc_super=1.0000
ACHIEVED PERFECT SCORE.

<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>

```

9.2 Sub Class Training

```

[ ]: if APPROACH == "two_models":
    # Subclass model
    model_sub = SingleHeadModel(num_classes=num_sub, head=SUPER_HEAD_TYPE).
    ↪to(device)
    criterion_sub = nn.CrossEntropyLoss()
    optimizer_sub = setup_backbone_training(
        model_sub,
        fine_tune_mode=FINE_TUNE_MODE,

```

```

        lr_full=LR,
        lr_head=LR_HEAD,
    )

run_sub = None
if USE_WANDB:
    run_sub = wandb.init(
        entity=WANDB_ENTITY,
        project=WANDB_PROJECT,
        name=make_run_name("two_models_sub_resnet"),
        group="two_models",
        config={
            "approach": "two_models",
            "head": "sub",
            "backbone": BACKBONE,
            "epochs": EPOCHS,
            "lr_full": LR,
            "lr_head": LR_HEAD,
            "fine_tune_mode": FINE_TUNE_MODE,
            "img_size": IMG_SIZE,
            'sub_head_type': SUB_HEAD_TYPE,
        },
    )

best_val_sub = 0.0

print("\nTraining subclass model:")
for epoch in range(1, EPOCHS + 1):
    train_metrics = train_one_epoch(
        model_sub,
        optimizer_sub,
        train_loader,
        criterion_sub,
        mode="single_head_sub",
    )
    val_metrics = eval_one_epoch(
        model_sub,
        val_loader,
        criterion_sub,
        mode="single_head_sub",
    )

    val_acc = val_metrics.get("val_acc_sub", 0.0)
    print(
        f"[Sub] Epoch {epoch}: "
        f"train_loss={train_metrics['loss']:.4f}, "
        f"val_loss={val_metrics['val_loss']:.4f}, "

```

```

        f"val_acc_sub={val_acc:.4f}"
    )

    if USE_WANDB:
        wandb.log(
            {
                "epoch": epoch,
                "sub_train_loss": train_metrics["loss"],
                "sub_val_loss": val_metrics["val_loss"],
                "sub_val_acc": val_acc,
            },
            step=epoch,
        )

    if val_acc > best_val_sub:
        best_val_sub = val_acc
        torch.save(
            model_sub.state_dict(),
            os.path.join(DATA_ROOT, "best_sub_model.pth"),
        )
        print(" Saved new best subclass model")

    if run_sub is not None:
        run_sub.finish()

```

```

<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>

```

Training subclass model:

```

[Sub] Epoch 1: train_loss=2.4963, val_loss=1.4422, val_acc_sub=0.7602
    Saved new best subclass model
[Sub] Epoch 2: train_loss=0.8537, val_loss=0.3808, val_acc_sub=0.9352
    Saved new best subclass model
[Sub] Epoch 3: train_loss=0.2495, val_loss=0.1471, val_acc_sub=0.9642
    Saved new best subclass model
[Sub] Epoch 4: train_loss=0.1251, val_loss=0.1204, val_acc_sub=0.9623
[Sub] Epoch 5: train_loss=0.0797, val_loss=0.1175, val_acc_sub=0.9613
[Sub] Epoch 6: train_loss=0.0569, val_loss=0.1131, val_acc_sub=0.9652
    Saved new best subclass model
[Sub] Epoch 7: train_loss=0.0529, val_loss=0.1042, val_acc_sub=0.9681
    Saved new best subclass model

```

```

[Sub] Epoch 8: train_loss=0.0372, val_loss=0.0824, val_acc_sub=0.9749
    Saved new best subclass model
[Sub] Epoch 9: train_loss=0.0328, val_loss=0.0974, val_acc_sub=0.9671
[Sub] Epoch 10: train_loss=0.0302, val_loss=0.1236, val_acc_sub=0.9652
[Sub] Epoch 11: train_loss=0.0179, val_loss=0.0716, val_acc_sub=0.9787
    Saved new best subclass model
[Sub] Epoch 12: train_loss=0.0262, val_loss=0.1177, val_acc_sub=0.9671
[Sub] Epoch 13: train_loss=0.0258, val_loss=0.0949, val_acc_sub=0.9691
[Sub] Epoch 14: train_loss=0.0225, val_loss=0.0767, val_acc_sub=0.9807
    Saved new best subclass model
[Sub] Epoch 15: train_loss=0.0185, val_loss=0.0988, val_acc_sub=0.9739

<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>

```

9.3 Eval

```

[ ]: TAU_SUB = 0.85

[ ]: if APPROACH == "two_models":
    best_super_path = os.path.join(DATA_ROOT, "best_super_model.pth")
    best_sub_path   = os.path.join(DATA_ROOT, "best_sub_model.pth")
    model_super.load_state_dict(torch.load(best_super_path, map_location=device))
    model_sub.load_state_dict(torch.load(best_sub_path, map_location=device))

    run_eval = None
    if USE_WANDB:
        run_eval = wandb.init(
            entity=WANDB_ENTITY,
            project=WANDB_PROJECT,
            name=make_run_name("two_models_eval_resnet"),
            group="two_models",
            config={
                "approach": "two_models",
                "head": "sub",
                "backbone": BACKBONE,
                "epochs": EPOCHS,
                "lr_full": LR,
                "lr_head": LR_HEAD,
                "fine_tune_mode": FINE_TUNE_MODE,
                "img_size": IMG_SIZE,
                'sub_head_type': SUB_HEAD_TYPE,
                'super_head_type': SUPER_HEAD_TYPE,
            }
        )

```

```

        'tau_sub': TAU_SUB,
        'tau_super': TAU_SUPER,
    },
)

analyze_tau_sub(model_sub, mode="sub_single_head")
analyze_tau_super(model_super, mode="super_single_head")
evaluate_on_val_with_novelty(model_super, mode="super_single_head",
                             tau_super=TAU_SUPER, tau_sub=TAU_SUB,
                             loader=val_loader, name="val_super_only")
evaluate_on_val_with_novelty(model_sub, mode="sub_single_head",
                             tau_super=TAU_SUPER, tau_sub=TAU_SUB,
                             loader=val_loader, name="val_sub_only")
evaluate_pseudo_novel_sub_with_novelty(
    model_sub, mode="sub_single_head", tau_sub=TAU_SUB
)
novelty_dashboard(model_super, mode="super_single_head",
                  tau_super=TAU_SUPER, tau_sub=TAU_SUB,
                  include_pseudo=False) # no pseudo-novel meaning for super
˓→head
novelty_dashboard(model_sub, mode="sub_single_head",
                  tau_super=TAU_SUPER, tau_sub=TAU_SUB,
                  include_pseudo=True)

if run_eval is not None:
    run_eval.finish()

else:
    print("APPROACH is not 'two_models'; skipping two-model training in this
˓→cell.")

```

```

<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>

```

Seen subclass max-prob stats:

```

count = 1034
mean  = 0.983
std   = 0.069
percentiles:
  p 1: 0.612
  p 5: 0.936

```

```
p10: 0.982  
p25: 0.996  
p50: 0.999  
p75: 1.000  
p90: 1.000  
p95: 1.000  
p99: 1.000
```

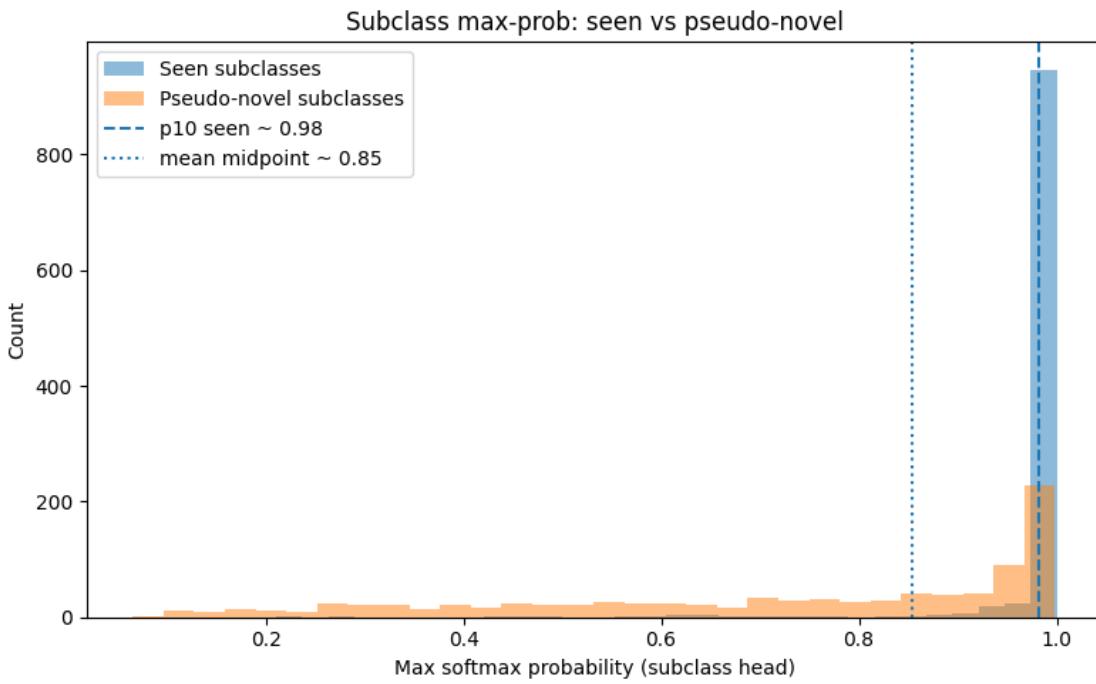
Pseudo-novel subclass max-prob stats:

```
count = 948  
mean = 0.724  
std = 0.264  
percentiles:  
p 1: 0.116  
p 5: 0.226  
p10: 0.307  
p25: 0.522  
p50: 0.810  
p75: 0.964  
p90: 0.993  
p95: 0.996  
p99: 0.997
```

Suggested TAU_SUB candidates:

```
tau_sub 10th percentile of seen: 0.982  
tau_sub mean(seen + pseudo)/2: 0.854
```

You can start with one of these for TAU_SUB and adjust based on leaderboard/behavior.



Logged histograms and tau_sub candidates to Weights & Biases.

Seen superclasses (0/1/2) superclass max-prob stats:

```

count = 534
mean  = 0.998
std   = 0.028
percentiles:
  p 1: 0.982
  p 5: 0.997
  p10: 0.999
  p25: 1.000
  p50: 1.000
  p75: 1.000
  p90: 1.000
  p95: 1.000
  p99: 1.000

```

Novel superclasses (== NOVEL_SUPER_IDX) superclass max-prob stats:

```

count = 500
mean  = 1.000
std   = 0.000
percentiles:
  p 1: 0.999
  p 5: 1.000
  p10: 1.000

```

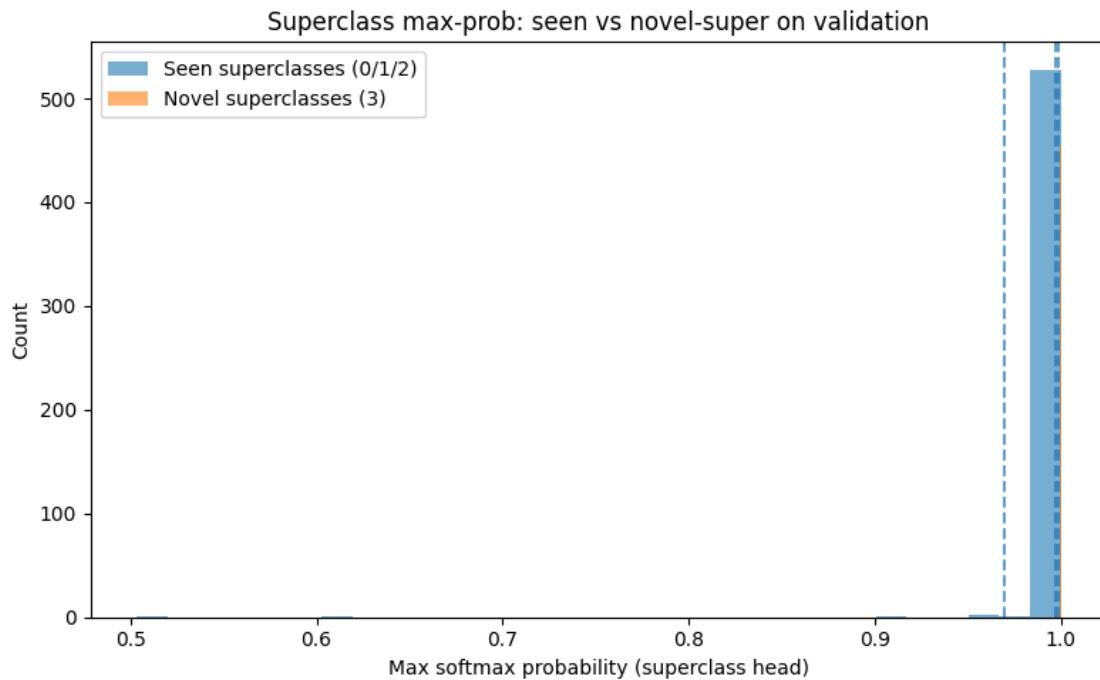
```

p25: 1.000
p50: 1.000
p75: 1.000
p90: 1.000
p95: 1.000
p99: 1.000

```

Suggested TAU_SUPER candidates (using seen + novel):

tau_super	10th percentile of seen:	0.999
tau_super	5th percentile of seen:	0.997
tau_super	mean(seen) - std(seen):	0.970
tau_super	mean(seen & novel) midpoint:	0.999



Logged superclass histograms and tau_super candidates to Weights & Biases.

```

==== Evaluation on val_super_only ====
Overall superclass acc: 0.9913
Seen superclass acc (true super != novel): 0.9831
Seen superclass false-novel rate: 0.0169
Novel superclass acc (true super == novel): 1.0000

==== Evaluation on val_sub_only ====
Overall subclass acc: 0.9584
Seen subclass acc (true sub != novel): 0.9195
Seen subclass false-novel rate: 0.0805

```

```
==== Evaluation on pseudo_novel_sub (held-out subclasses) ====
Fraction flagged as novel (good): 0.5464
Fraction mapped to seen subclasses (bad): 0.4536
```

```
==== Evaluation on val ===
Overall superclass acc: 0.9913
Seen superclass acc (true super != novel): 0.9831
Seen superclass false-novel rate: 0.0169
Novel superclass acc (true super == novel): 1.0000
```

```
===== Novelty Dashboard =====
```

	Split	Head	Metric \
0	config	-	APPROACH
1	config	-	BACKBONE
2	config	-	CIFAR_NOVEL_MODE
3	config	-	DATA_AUGMENT
4	config	-	FINE_TUNE_MODE
5	config	-	TAU_SUB
6	config	-	TAU_SUPER
7	config	-	USE_PSEUDO_NOVEL
8	val	super	Novel superclass accuracy (CIFAR)
9	val	super	Seen superclass accuracy
10	val	super	Seen superclass false-novel rate

	Meaning	Value
0	Model architecture (two_heads vs two_models)	two_models
1	Feature extractor (e.g. resnet18 / resnet50)	resnet50
2	Extra novel-super CIFAR data mode	large
3	Whether data augmentation is enabled for training	True
4	Backbone training mode (full vs frozen)	full
5	Novelty threshold for subclass head	0.85
6	Novelty threshold for superclass head	0.99
7	Using held-out subclasses for pseudo-novel eval	True
8	CIFAR novel-super samples correctly predicted ...	1.0
9	Correctly keep seen superclasses as seen	0.983146
10	Seen superclasses incorrectly flipped to novel	0.016854

```
==== Evaluation on val ===
```

```
Overall subclass acc: 0.9584
Seen subclass acc (true sub != novel): 0.9195
Seen subclass false-novel rate: 0.0805
```

```
==== Evaluation on pseudo_novel_sub (held-out subclasses) ====
Fraction flagged as novel (good): 0.5464
Fraction mapped to seen subclasses (bad): 0.4536
```

```
===== Novelty Dashboard =====
```

```

          Split Head                                Metric \
0      config    -                               APPROACH
1      config    -                               BACKBONE
2      config    -           CIFAR_NOVEL_MODE
3      config    -           DATA_AUGMENT
4      config    -           FINE_TUNE_MODE
5      config    -           TAU_SUB
6      config    -           TAU_SUPER
7      config    -           USE_PSEUDO_NOVEL
8  pseudo_novel sub   Pseudo-novel mapped to seen
9  pseudo_novel sub   Pseudo-novel marked as novel
10         val sub   Seen subclass accuracy
11         val sub   Seen subclass false-novel rate

                                         Meaning          Value
0      Model architecture (two_heads vs two_models) two_models
1      Feature extractor (e.g. resnet18 / resnet50) resnet50
2                  Extra novel-super CIFAR data mode large
3  Whether data augmentation is enabled for training True
4                  Backbone training mode (full vs frozen) full
5                  Novelty threshold for subclass head 0.85
6                  Novelty threshold for superclass head 0.99
7  Using held-out subclasses for pseudo-novel eval True
8  Held-out subclasses wrongly mapped to some see... 0.453586
9  Held-out subclasses correctly flagged as novel 0.546414
10     Correctly keep seen subclasses as seen 0.919476
11     Seen subclasses incorrectly flipped to novel 0.080524

<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
```

10 SECTION 8: Test-time Inference & CSV Export for leader-board (two-head model)

```
[ ]: # Inference in order to predict "novel" superclass / subclasses
```

```

@torch.no_grad()
def predict_test_two_heads(model, test_loader, tau_super=TAU_SUPER, ↴
    ↴tau_sub=TAU_SUB):
    """
    Inference for the two-heads model with novelty thresholds.

```

```

model:      SharedBackboneTwoHeads(...)
test_loader: DataLoader yielding (images, img_names)
tau_super:  threshold for superclass novelty
tau_sub:   threshold for subclass novelty
"""

model.eval()
images_list = []
super_preds = []
sub_preds = []

for images, img_names in test_loader:
    images = images.to(device, non_blocking=True)

    # Forward pass once
    super_logits, sub_logits = model(images)      # (B, num_super), (B, num_sub)

    # --- Superclass predictions with novelty ---
    super_probs = F.softmax(super_logits, dim=1)           # (B, num_super)
    max_super_probs, super_idx = super_probs.max(dim=1)      # (B, )
    super_novel_mask = max_super_probs < tau_super
    super_idx = super_idx.clone()
    super_idx[super_novel_mask] = NOVEL_SUPER_IDX

    # --- Subclass predictions with novelty ---
    sub_probs = F.softmax(sub_logits, dim=1)                # (B, num_sub)
    max_sub_probs, sub_idx = sub_probs.max(dim=1)            # (B, )
    sub_novel_mask = max_sub_probs < tau_sub
    sub_idx = sub_idx.clone()
    sub_idx[sub_novel_mask] = NOVEL_SUB_IDX

    # Move to CPU as plain Python ints
    super_idx = super_idx.cpu().tolist()
    sub_idx = sub_idx.cpu().tolist()

    # img_names is a list of filenames (len = B)
    images_list.extend(img_names)
    super_preds.extend(super_idx)
    sub_preds.extend(sub_idx)

df = pd.DataFrame({
    "image": images_list,
    "superclass_index": super_preds,
    "subclass_index": sub_preds
})
return df

```

```

# SECTION: Test-time inference & CSV export for two-heads model

if APPROACH == "two_heads":
    # Recreate the model and load best checkpoint
    model_two_heads = SharedBackboneTwoHeads(num_super=num_super, num_sub=num_sub).to(device)
    best_ckpt_path = os.path.join(DATA_ROOT, "best_two_heads_kl.pth")
    model_two_heads.load_state_dict(torch.load(best_ckpt_path, map_location=device))

    test_predictions = predict_test_two_heads(
        model_two_heads,
        test_loader,
        tau_super=TAU_SUPER,
        tau_sub=TAU_SUB,
    )
    out_csv_path = os.path.join(DATA_ROOT, "two_heads_predictions.csv")
    test_predictions.to_csv(out_csv_path, index=False)
    print("Saved two-heads predictions (with novelty) to:", out_csv_path)
else:
    print("APPROACH is not 'two_heads'; skipping two-heads inference in this cell.")

```

APPROACH is not 'two_heads'; skipping two-heads inference in this cell.

11 SECTION 9: Test-time Inference & CSV Export for leader-board (two separate models)

[]: TAU_SUB

[]: 0.85

```

[ ]: @torch.no_grad()
def predict_test_two_models(model_super, model_sub, test_loader,
                           tau_super=TAU_SUPER, tau_sub=TAU_SUB):
    """
    Inference for the two-model setup (separate super + sub models) with novelty thresholds.

    model_super: SingleHeadModel for superclass (num_classes = num_super)
    model_sub: SingleHeadModel for subclass (num_classes = num_sub)
    test_loader: DataLoader yielding (images, img_names)
    tau_super: threshold for superclass novelty
    tau_sub: threshold for subclass novelty
    """

```

```

model_super.eval()
model_sub.eval()

images_list = []
super_preds = []
sub_preds = []

for images, img_names in test_loader:
    images = images.to(device, non_blocking=True)

    # Forward passes
    super_logits = model_super(images)      # (B, num_super)
    sub_logits   = model_sub(images)        # (B, num_sub)

    # --- Superclass predictions with novelty ---
    super_probs = F.softmax(super_logits, dim=1)           # (B, num_super)
    max_super_probs, super_idx = super_probs.max(dim=1)     # (B, )
    super_novel_mask = max_super_probs < tau_super
    super_idx = super_idx.clone()
    super_idx[super_novel_mask] = NOVEL_SUPER_IDX

    # --- Subclass predictions with novelty ---
    sub_probs = F.softmax(sub_logits, dim=1)                 # (B, num_sub)
    max_sub_probs, sub_idx = sub_probs.max(dim=1)            # (B, )
    sub_novel_mask = max_sub_probs < tau_sub
    sub_idx = sub_idx.clone()
    sub_idx[sub_novel_mask] = NOVEL_SUB_IDX

    # Move indices to CPU as plain Python ints
    super_idx = super_idx.cpu().tolist()
    sub_idx   = sub_idx.cpu().tolist()

    # img_names is a list of filenames (len = B)
    images_list.extend(img_names)
    super_preds.extend(super_idx)
    sub_preds.extend(sub_idx)

df = pd.DataFrame({
    "image": images_list,
    "superclass_index": super_preds,
    "subclass_index": sub_preds,
})
return df

if APPROACH == "two_models":

```

```

    model_super = SingleHeadModel(num_classes=num_super, head=SUPER_HEAD_TYPE) .
    ↪to(device)
    model_sub   = SingleHeadModel(num_classes=num_sub, head=SUB_HEAD_TYPE) .
    ↪to(device)

    model_super.load_state_dict(torch.load(os.path.join(DATA_ROOT, "best_super_model.pth"),
    ↪map_location=device))
    model_sub.load_state_dict(torch.load(os.path.join(DATA_ROOT, "best_sub_model.pth"),
    ↪map_location=device))

    test_predictions_two_models = predict_test_two_models(
        model_super, model_sub, test_loader,
        tau_super=TAU_SUPER, tau_sub=TAU_SUB,
    )

    out_csv_path = os.path.join(DATA_ROOT, "two_models_predictions.csv")
    test_predictions_two_models.to_csv(out_csv_path, index=False)
    print("Saved two-model predictions (with novelty) to:", out_csv_path)

else:
    print("APPROACH is not 'two_models'; skipping two-models inference in this
    ↪cell.")

```

Saved two-model predictions (with novelty) to: /content/drive/MyDrive/NNDL-Project/Project Data/two_models_predictions.csv

11.1 Submission Info

Team Name: Caged Manifolds

Model Name: ResNet50v2 with Data Augmentation - Cosine Similarity heads

Description: 2 model architecture (separate super and subclass predictor), various regularization and data augmentation added, cosine classifier subhead