
Transfer Learning: Exploring Two-Headed CNN w/Interaction Loss vs. Two separate CNN Models for Superclass/Subclass Classification

Christian Montgomery and Aksel Kretsinger-Walters

Columbia University

cm4521@columbia.edu, adk2164@columbia.edu

1 Introduction

We explored optimal techniques of transfer learning for the multi-label image classification task, predicting superclass (bird/dog/reptile/novel class) and subclasses (87 seen + novel subclass) for a selection of 64×64 images. The main challenge to overcome was distribution shift between the training and test datasets, and novel classes appearing at test time (identifying both new superclasses and new subclasses as “novel”).

To maximize our performance on this classification, we explored two different model architectures: a multi-headed CNN with KL Interaction term, and separately two distinct CNN’s for each classification task. We trained detecting for “novel” superclasses via including some out-of-distribution data from CIFAR 100, and held-out some subclass images from training to assess pseudo-novel detection. For both classifications, we tuned a threshold on the max softmax score for each image, implying whether it represented a novel class or not. Two other areas we explored were varying the amount of out of distribution “novel” superclass images provided during training and experimenting with both full and frozen fine-tuning to assess generalization on unseen classes during testing.

Overall, testing showed that two-separate models performed better at identifying seen superclasses, worse at unseen superclasses, and about equivalent on subclasses overall. Interestingly, reducing the number of CIFAR 100 novel superclass images included (from 5,000 to 1,000) did not alter superclass prediction much, contrary to expectations. But it did meaningfully worsen performance on seen subclasses while meaningfully improving performance on unseen subclasses (perhaps expected from a model more “naive” to distribution of natural images). Lastly, altering the fine tuning from full to frozen changed.....

2 Background & Related Work

Our work sits at the intersection of transfer learning for vision, multi-task classification and novelty detection. A common approach in computer vision is to use a CNN pretrained on a large dataset like ImageNet and then adapt it for a downstream task. Our approach considered both ends of the fine-tuning spectrum, 1. full fine-tuning of all ResNet layers which has been shown to exhibit high performance at the potential cost of overfitting, and 2. freezing the backbone to extract features from the base model while training only the linear head on top which is known to lead to greater stability.

Prior work on multi-task and multi-label classification often uses one of two architectures, either a shared feature backbone with multiple heads, or multiple models. We decided to implement both patterns, as the shared-backbone with multiple heads is able to share features between the heads, while the two-separate model design allows each classifier to specialize. builds upon previous transfer learning experimentation and has discovered that the combination of X, Y, and Z provides the best performance for this specific task.

A key challenge is that our test data includes classes that aren’t available in training. This relates to the concept of open-set recognition and out-of-distribution (OOD) detection, where the goal is to flag samples that don’t belong to any known training class. A common approach is to apply a threshold on the maximum softmax probability, and if the model’s confidence is below this threshold, the sample is treated as “novel”, which is what we use. In order to train our model to recognize images from the subspace of “novel” superclasses, we incorporate OOD training by including images from CIFAR-100, to approximate the distribution of unknown classes. For novel subclasses, without a curated set of new dog/bird/reptile subclasses, we used a common “pseudo-novel” strategy of holding out a subset of subclasses from the training set, which allowed us to calibrate subclass thresholds.

Lastly, in the case of a single model with a shared backbone and two separate heads, it is common to use an interaction loss (KL divergence) to encourage consistency between superclass and subclass predictions (since each subclass is mapped to a unique

superclass). This type of consistency regularization relates to ideas from knowledge distillation and hierarchical classification, where one head is encouraged to agree with a derived distribution from another.

3 Methods

Overall, we built our image classification model upon a CNN-based architecture, using PyTorch within a Colab GPU and Weights & Biases for tracking training and validation error. We chose a convolutional neural network based architecture as they have repeatedly been shown to perform well on image classification tasks. We leveraged a pretrained ResNet-50 backbone trained on ImageNet to transfer learn from, and added a linear layer on top to map from final feature vectors of the pretrained model to our classification labels (both superclass and subclass). There were about $\sim 6,600$ training images, 64×64 RGB, which had superclass labels of {bird, dog, reptile} for training, 87 subclass labels within each superclass, and a “novel” label used for both unseen superclasses and unseen subclasses at test.

In processing our input data, we resized if necessary to 64×64 and then normalized with standard ImageNet normalization using the average and standard deviation of pixel values used for the RGB channels of ImageNet dataset. We randomly applied horizontal flips and random rotations by 10 degrees to our training data in order to ensure our model learned image content while remaining invariant to subject orientation.

3.1 CIFAR-100 based novel superclasses

To ensure our model was trained to generalize to images with a superclass it has never seen before (and thus should label novel), we supplemented with a subset of CIFAR-100 images which represented a novel superclass (and by extension subclass). For example, these were images that did not represent a bird, dog, or reptile, and thus we labeled them as novel. We started with 5,000 additional images (roughly double training set size) and experimented with scaling that down to only 1,000 additional images ($\sim 15\%$ of training set size). We had our validation set as 10% of total training data.

3.2 Pseudo-novel subclasses

For our model to generalize to novel subclasses, unfortunately we were limited by time and resources and could not curate a precise dataset of dogs, reptiles and birds that were exclusively not of the subclass categories provided in order to train on those novel subclasses. Instead, we used the concept of “pseudo-novel” subclasses. This entailed randomly holding out a fraction (15%) of subclasses from training entirely, then training on “seen” subclasses only, and using held-out subclasses during validation to determine an optimal confidence threshold by which to classify an image as a novel subclass. For inference, we did not hold out any subclasses so the model would have full knowledge of all subclasses in testing.

3.3 Thresholding rule

For a mechanism to infer whether a test image was of a superclass and subclass previously seen, we introduced a thresholding rule for both classification. This means that for a given percentage, (i.e. 80%), if the max softmax score for a given classification task is above 80%, the class with that softmax score is chosen as the predicted class, but if score is below, then “novel” class is chosen instead. We used separate τ_{super} and τ_{sub} for thresholding the superclass and subclass predictions respectively, which were tuned by analyzing validation performance. The below equation represents how that calculation was used at inference for both classification tasks.

$$\hat{y} = \begin{cases} \arg \max_c p(c | x), & \text{if } \max_c p(c | x) \geq \tau \\ \text{novel}, & \text{otherwise.} \end{cases}$$

Before running our model for inference, we tuned these thresholds by toggling on “pseudo-novel” so that a subset of subclasses were held out from training. We charted the results below, and decided on $\tau_{\text{super}} = 0.99$ and $\tau_{\text{sub}} = 0.85$ accordingly. The tradeoff was needing to balance between a lower threshold meant missing some “Novel” classifications, while a higher threshold would mean incorrectly classify some seen classes as “Novel”. The τ_{super} of 0.99 though high, makes sense, as we augmented training data with “Novel” CIFAR-100 images so our model actually learns about the “Novel” superclass subspace.

3.4 Approach A: Single model, 2 separate heads + KL divergence term

Our first approach in Model Architecture was utilizing a shared ResNet-50 backbone for feature learning and having two classification heads (superclass and subclass). The superclass head produced superclass logits (4 classes including novel) and the subclass head produced subclass logits. Each head is trained using its own cross-entropy loss error term, and we added an additional term to mitigate potential divergence between superclass and subclass predictions (as in, the superclass being a bird and the subclass being a golden retriever). Accordingly, we added an interaction loss term, which is the KL divergence

between: 1. Implied probability of a superclass (sum of corresponding subclass probabilities) and 2. The superclass model head's calculated probability of that subclass. Thus, the total loss reflects this term as well.

Diagram below shows our total loss computation by component.

3.5 Model Architecture

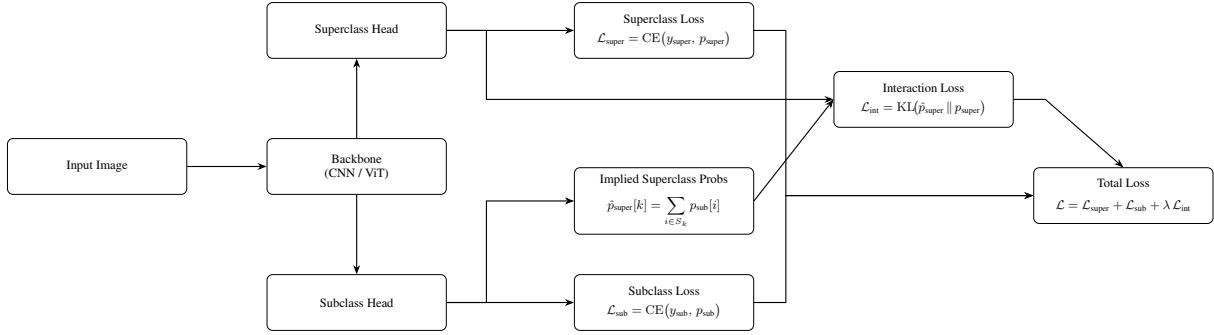


Figure 1: Model architecture and loss computation.

3.6 Interaction Loss / KL Divergence Calculations

For each superclass k , with subclass set S_k :

$$\hat{p}_{\text{super}}[k] = \sum_{i \in S_k} p_{\text{sub}}[i].$$

Then define interaction loss, e.g.:

$$\mathcal{L}_{\text{int}} = \text{KL}(\hat{p}_{\text{super}} \parallel p_{\text{super}}).$$

The Kullback–Leibler (KL) divergence between two distributions P and Q is:

$$\text{KL}(P \parallel Q) = \sum_k P(k) \log \frac{P(k)}{Q(k)}.$$

Equation below demonstrates math behind the KL divergence, which we used in Approach A.

3.7 Approach B: Two independent models, separate for superclass and subclass

Our second approach was adapting our architecture to have two separate ResNet-50 models, one for superclass and one for subclass. Each model has its own cross-entropy loss during training, and they are independently trained. At inference, their predictions are combined. The benefits of this approach are that each model can specialize for its particular classification task with little inference, the downside is more parameters and no consistency regularization.

3.8 Approach C: Two independent models, separate for superclass and subclass (cosine similarity head)

Our third approach was to modify the subclass model to implement a cosine similarity head instead of a linear feed-forward head.

3.9 Evaluation

The metrics used to evaluate model performance on unlabeled test data were: 1. cross-entropy, 2. overall accuracy, 3. seen accuracy and 4. unseen accuracy, all for both superclass and subclass. We were limited to a max of five leaderboard submissions.

4 Experimental Setup

Our hyperparameters were chosen as the following (for all models, unless denoted above): Batch size of 64, 15 training epochs, learning rate of 10^{-4} , weight decay of 10^{-4} for generalization, and Adam as optimizer. The hyperparameters used for the KL

divergence term in Approach A was an alpha value of 0.1 and a Temperature value of 1.0. Lastly, We selected a τ_{super} of 0.99 and τ_{sub} of 0.85 to threshold novel class predictions.

Below is a configuration table of our submissions, to illustrate differences in approaches.

Table 1: Configuration table of submissions.

Approach	Architecture	CIFAR novel super size	Backbone	Learning Rate	Head Type
A	Two-heads	5000	fine-tuned	.0001	Linear
B	Two models	5000	fine-tuned	.0001	Linear
C	Two models	5000	fine-tuned	.0001	Super-Linear, Sub-Linear

5 Results

Below table shows overall performance vs. baseline - the leaderboard test accuracy metrics:

Table 2: Leaderboard test accuracy metrics.

Approach	Overall Super	Seen Super	Unseen Super	Overall Sub	Seen Sub	Unseen Sub
A	83%	90%	63%	71%	84%	67%
B	85%	96%	57%	70%	83%	66%
C	90%	94%	79%	70%	75%	68%

5.1 Analysis

Overall, right from our initial Approach A and leaderboard submission we saw improved test accuracy vs. baseline in all categories except “Seen Super” where we had 90% accuracy for Approach A vs. 99% baseline. This made sense to us, as there is an inherent tradeoff in maximizing seen vs. unseen accuracies when using a threshold to predict “Novel” classes. The tradeoff is that if you set a lower threshold, say 90% threshold instead of 99% threshold as we did, we may have improved “Seen” accuracy at the expense of “Novel” accuracy since we would have reduced the rate of misclassifications labeled as “Novel” that were actually “Seen”. Since the goal was to improve the unseen accuracies vs. baseline, which we did meaningfully for both Super and Sub, we are satisfied with these results. Thus, we kept these thresholds (τ_{sub} and τ_{super}) consistent for all of our approaches accordingly. We will now review the differences each approach was meant to isolate, one by one.

5.2 Two-Heads vs. Two-Models performance

Below demonstrates the key performance differences between these two approaches, with all else held constant. Notably, only the superclass performance diverged meaningfully, as subclass performance was very similar. The superclass performance on “Seen” superclasses improved for two-models, which is intuitive as there was no interaction between the two classifier heads (as the two-headed model faced). However, the generalization worsened as the superclass performance on “Unseen” superclasses decreased for two-models, which makes sense as the superclass model became relatively more overfit without the implicit regularization of the shared backbone & KL-term.

Performance Delta Between Two-Heads and Two-Models

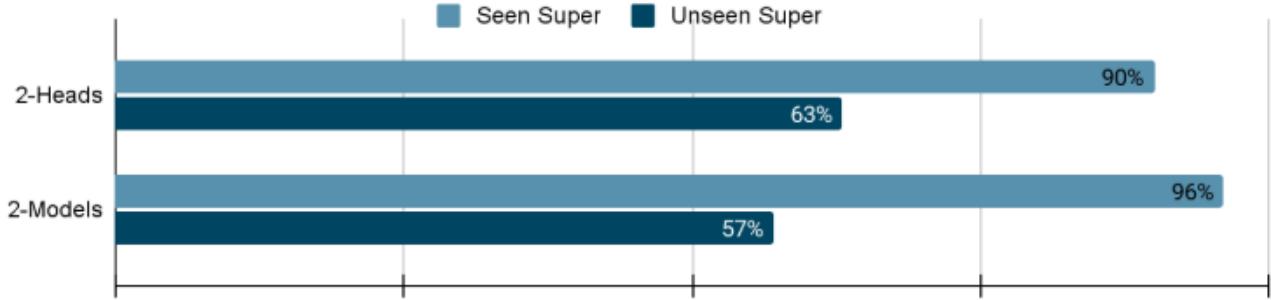


Figure 2: Performance delta between Two-Heads and Two-Models.

5.3 5,000 vs. 1,000 CIFAR-100 “Novel” Superclass Images Added performance

Interestingly, reducing the number of “Novel” class superclass images added to the training data only affected the superclass accuracy performance very little. We suppose this is from 1,000 CIFAR-100 being of sufficient size (relative to $\sim 6,600$ original training images) for the model to learn the patterns of the subspace of images that do not represent a bird, dog, or reptile, such that additional random examples from this complement space do not notably improve the accuracy.

However, for subclass accuracy, we observe meaningful differences. Notably, when we reduce the amount of images added (labeled as Novel superclass), we worsen performance on seen subclasses, but meaningfully improve performance on unseen subclasses. Given that this was tested on the two-model architecture, where the only training data with “Novel” subclasses would be images that are not birds, dogs and reptiles, the most plausible explanation is that we were “teaching” the model that novel subclasses come from a distribution beyond the superclasses of the training data. Therefore, when we relax this constraint (less images added), we reduce the “overfitting” to that constrained case of subclasses out of training distribution, and thus we improve on generalization of subclass accuracy for unseen subclasses in test data. We expect the model with less images added is more “naive” to the distribution of natural images.

Performance Delta Between 5,000 and 1,000 Superclass Images Added

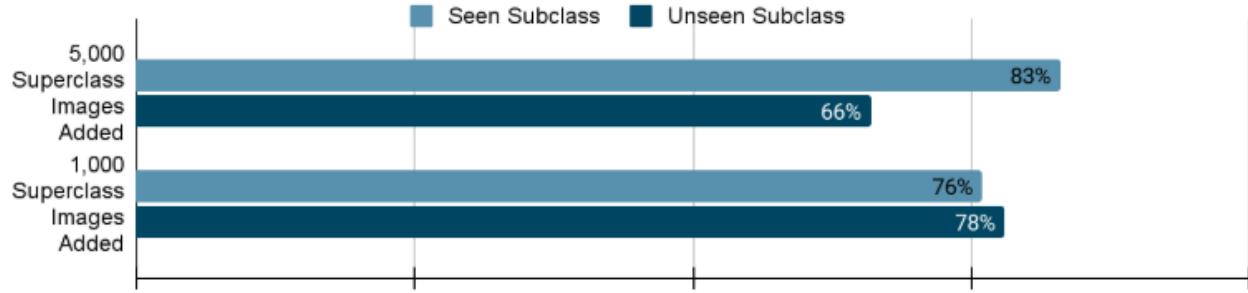


Figure 3: Performance delta between 5,000 and 1,000 superclass images added.

5.4 Amount of Fine-Tuning for Transfer Learning (Full fine-tuning vs. Fixed backbone)

Here we changed from full fine-tuning (with a learning rate of 0.0001) to fixing the ResNet50 backbone and only fine-tuning a single linear layer on top (with a learning rate of 0.01). This led to our worst performance, with the performance on subclasses meaningfully worse across the board vs. any of our other approaches (for overall, seen, and super). This makes sense as we likely had too high a learning rate to learn nuances between dogs/reptiles/birds well. The results for superclass accuracy were more interesting, as the seen superclass performance was worse, but the unseen superclass performance the best out of all of our approaches. We think this is because only the single linear layer was tuned on the superclass model, so it was able to specialize in identifying a “Novel” superclass vs. having that signal distorted when propagated across the full backbone (as in our other approaches).

Performance Delta Between Full Fine-Tuning and Fixed Backbone

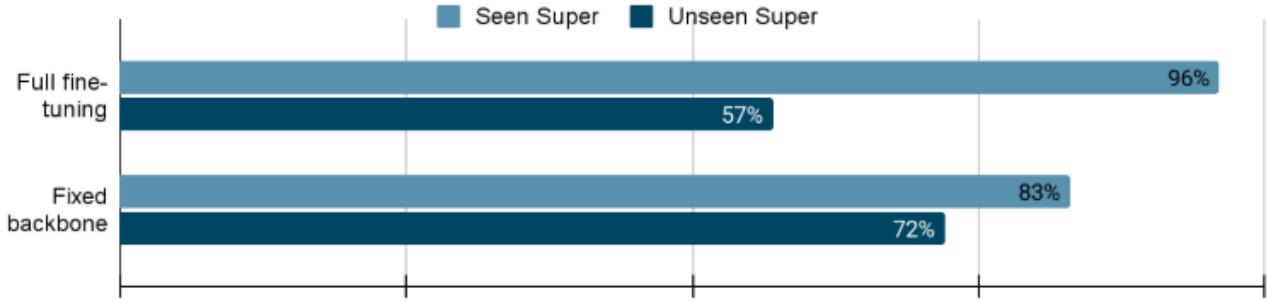


Figure 4: Performance delta between full fine-tuning and fixed backbone.

6 Discussion and Limitations

Our best model demonstrates (Approach B) demonstrates good performance on seen super and subclasses, and $\sim 60\%$ performance of catching novel super and subclasses in test data. We note that at the time of writing, this model’s performance seems to exceed that of most other ResNet-based architectures on the leaderboard, though surpassed by vision transformer-based architectures.

Limitations exist in our use of CIFAR and “pseudo-novel” subclass holdout as proxies for calibrating our thresholds for inference of “Novel” superclass and subclass respectively. CIFAR-100 served as an approximation for showing the model “out of distribution” superclass data, but did not necessarily include ambiguous data on the margin. Neither did our “pseudo-novel” technique necessarily teach our model to learn the full subspace of dog breeds (or reptiles/birds). Since there is a natural confidence overlap between seen and novel, we always faced the tradeoff in determining at what discrete point the boundary should lie.

We initially expected a single model with two heads and KL divergence term would be higher performing, due to its seemingly greater sophistication and unifying approach. We were surprised to learn, however, that two models trained separately perform better on our test data. Additionally, our initial instinct was that freezing the transfer learning backbone and only fine-tuning a layer on top would be superior, but apart from the unseen superclass evaluation metric, it was not. Lastly, it was also very surprising that reducing the amount of CIFAR superclass data affected the subclass performance much more noticeably than the superclass performance, though now we better appreciate the inherent complexity and sensitivity to change of neural networks.

7 Conclusion & Future Work

We experimented with various strategies to maximize performance on this transfer-learning image classification task. The approaches pursued were changing architecture (single model with two heads + KL divergence term vs. two models), changing number of extra images added to training data (5,000 vs. 1,000), and changing the amount of fine-tuning (full fine-tuning vs. fixed backbone). Our best model was using two separate models, with 5,000 novel superclass images added, and performing full fine-tuning.

Given further time and resources, we would have spent more time curating out of distribution datasets to add to the training data (in particular, images from novel subclasses from the given superclasses, as in new dog breeds, reptiles, and birds). We would have explored implementing more advanced and modern architecture such as Vision-Transformers.

References

- [1] **J. Donahue, Y. Jia, O. Vinyals, J. Hoffman, N. Zhang, E. Tzeng, and T. Darrell.** DeCAF: A Deep Convolutional Activation Feature for Generic Visual Recognition. In *Proceedings of the 31st ICML*, 2014.
- [2] **A. Krizhevsky, I. Sutskever, and G. E. Hinton.** ImageNet Classification with Deep Convolutional Neural Networks. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2012.
- [3] **R. Caruana.** Multitask Learning. *Machine Learning*, 28(1):41–75, 1997.
- [4] **G. Hinton, O. Vinyals, and J. Dean.** Distilling the Knowledge in a Neural Network. In *NIPS Deep Learning and Representation Learning Workshop*, 2015.
- [5] **K. He, X. Zhang, S. Ren, and J. Sun.** Deep Residual Learning for Image Recognition. In *CVPR*, 2016.
- [6] **J. Yosinski, J. Clune, Y. Bengio, and H. Lipson.** How Transferable Are Features in Deep Neural Networks? In *NeurIPS*, 2014.
- [7] **S. J. Pan and Q. Yang.** A Survey on Transfer Learning. *IEEE Transactions on Knowledge and Data Engineering*, 22(10):1345–1359, 2010.
- [8] **A. Bendale and T. E. Boult.** Towards Open Set Deep Networks. In *CVPR*, 2016.
- [9] **D. Hendrycks and K. Gimpel.** A Baseline for Detecting Misclassified and Out-of-Distribution Examples in Neural Networks. In *ICLR*, 2017.
- [10] **P. Goyal, P. Dollár, R. Girshick, P. Noordhuis, L. Wesolowski, A. Kyrola, A. Tulloch, Y. Jia, and K. He.** Accurate, Large Minibatch SGD: Training ImageNet in 1 Hour. *arXiv:1706.02677*, 2017.